

# MONTE CARLO TREE SEARCH



A heuristic search algorithm for some kinds of decision processes, most notably those employed in software that plays board games.

SUPERVISED BY:

**DR. / Mohamed Hussien**

# Team Members

1. Ali Mohamed Ali
2. Omar Ahmed Abd-Alaziz
3. Marios Magid Atef
4. Hagar Ali El-Refaai
5. Hana Ahmed Nabhan



# AGENDA

- Introduction
- Principle of operation
- The 4 Phases of MCTS
  - 1. Selecting
  - 2. Expansion
  - 3. Simulation
  - 4. Backpropagation
  - Exploration and Exploitation
  - UCT Formula
- Flowchart
- Pseudocode
- Worked Example
  - 1. Initial state
  - 2. First iteration
  - 3. second iteration
  - 4. Third iteration
- Advantages & Disadvantages
- Common Q&A



# Introduction

## Definition

Monte Carlo Tree Search (MCTS) is a **simulation-based tree search algorithm** that uses random sampling to make decisions. It is widely used in game AI (such as Go, Chess, and Tic-Tac-Toe) and in complex planning problems.

## Key Idea:

Build a search tree incrementally using **random simulations** to estimate the value of moves



# Introduction

## Where It's Used:

Used in game play.

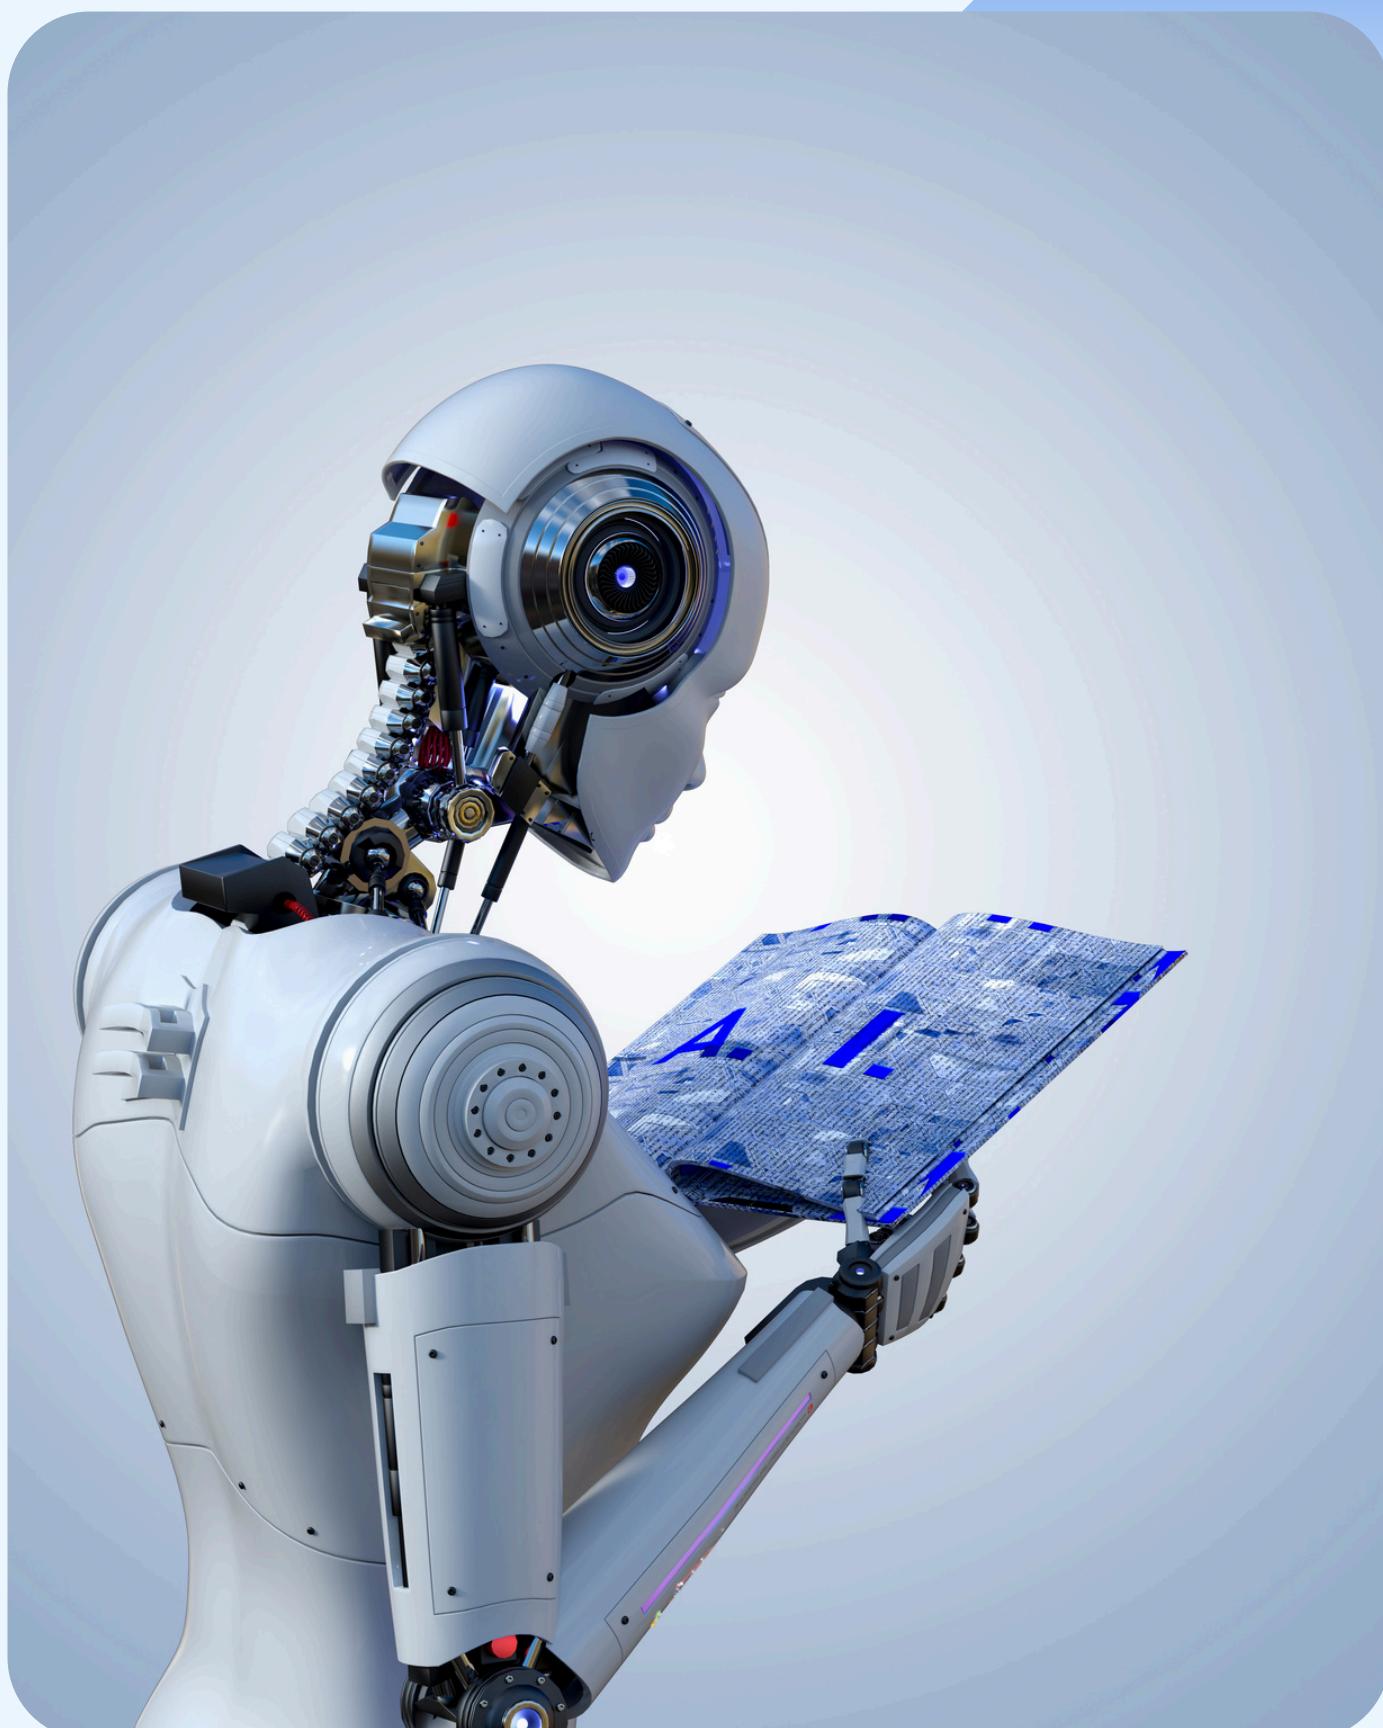
- o MCTS was introduced in **2006** for computer Go
- o Used in **board games** like **chess** and **shogi**
- o Used in **games with incomplete information** such as **bridge** and **poker**
- o Used in **real-time video games** such as Total War: Rome II's implementation of the high level campaign A1.





# Principle of operation

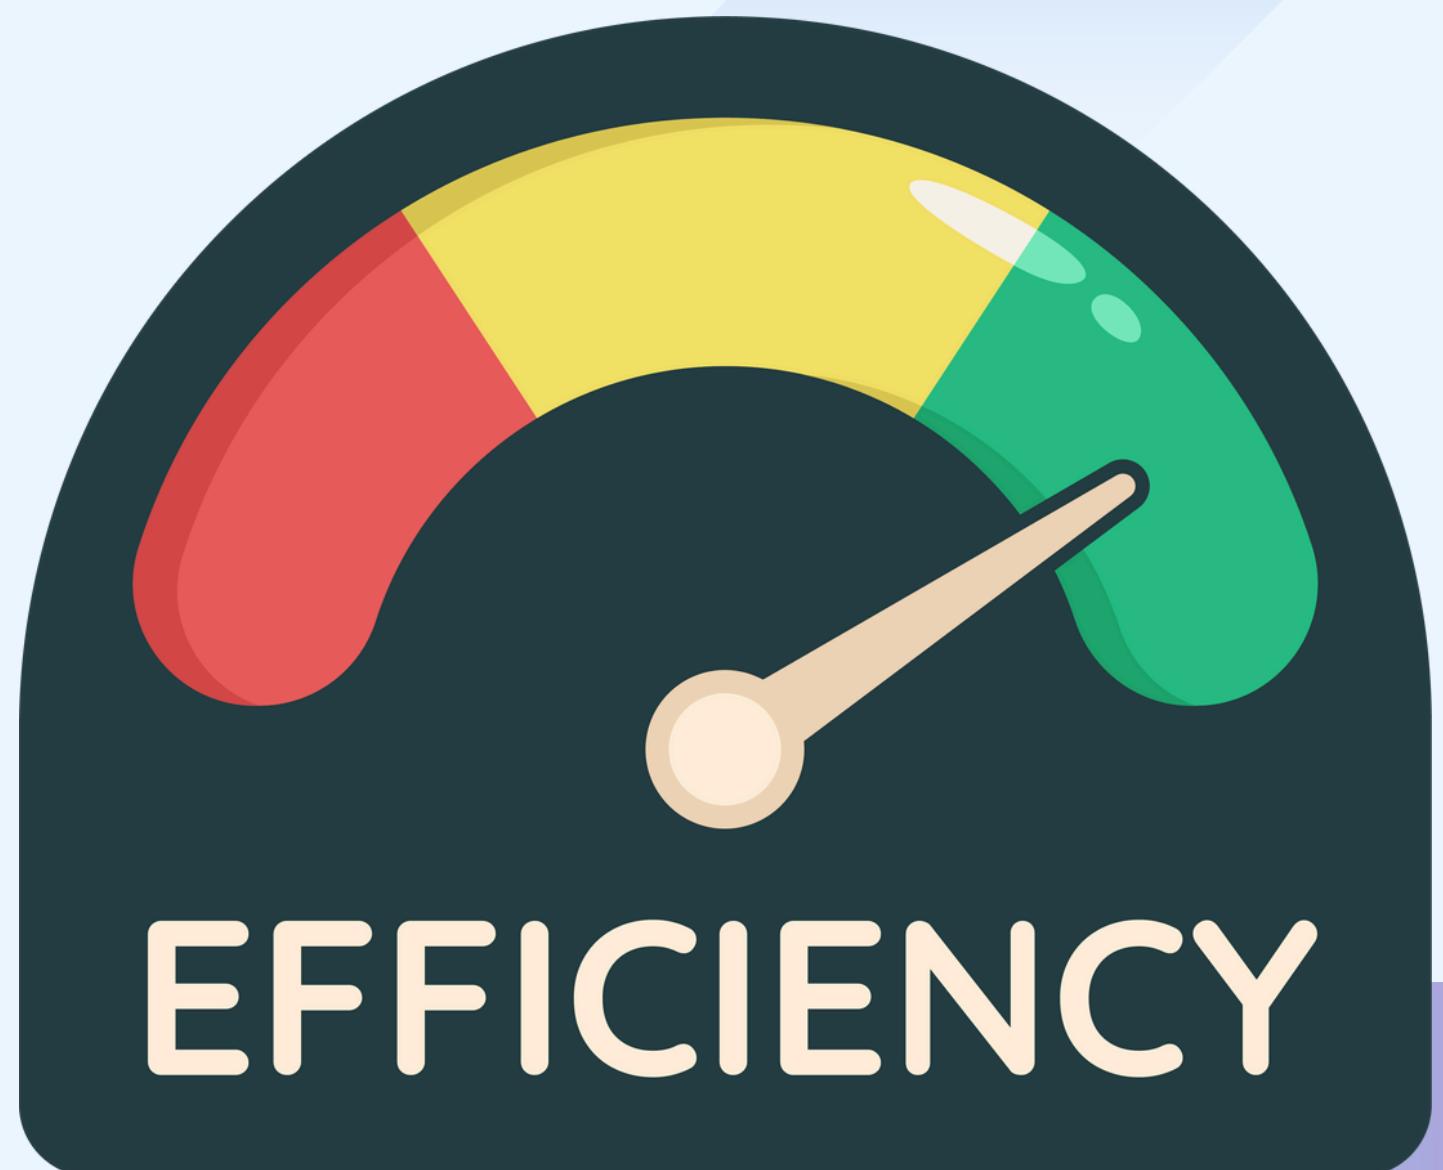
- The focus of MCTS is on **the analysis** of the most **promising moves**.
- It expands the search tree based on **random sampling** of the search space.
- The application of Monte Carlo tree search in games is based on many **playouts**, also called roll-outs.
- In each playout, the game is played out to the very end by **selecting moves at random**.





# Principle of operation

- The final game result of each playout is then used to **weight the nodes** in the game tree so that better nodes are more likely to be chosen in future playouts.
- The **efficiency** of this method—called **Pure Monte Carlo Game Search**—often **increases** with time as more playouts are assigned to the moves that have frequently resulted in the current player's victory according to previous playouts.



# The 4 Phases of MCTS

Each round of Monte Carlo tree search consists of four steps

## Selecting



## Expansion



## Simulation

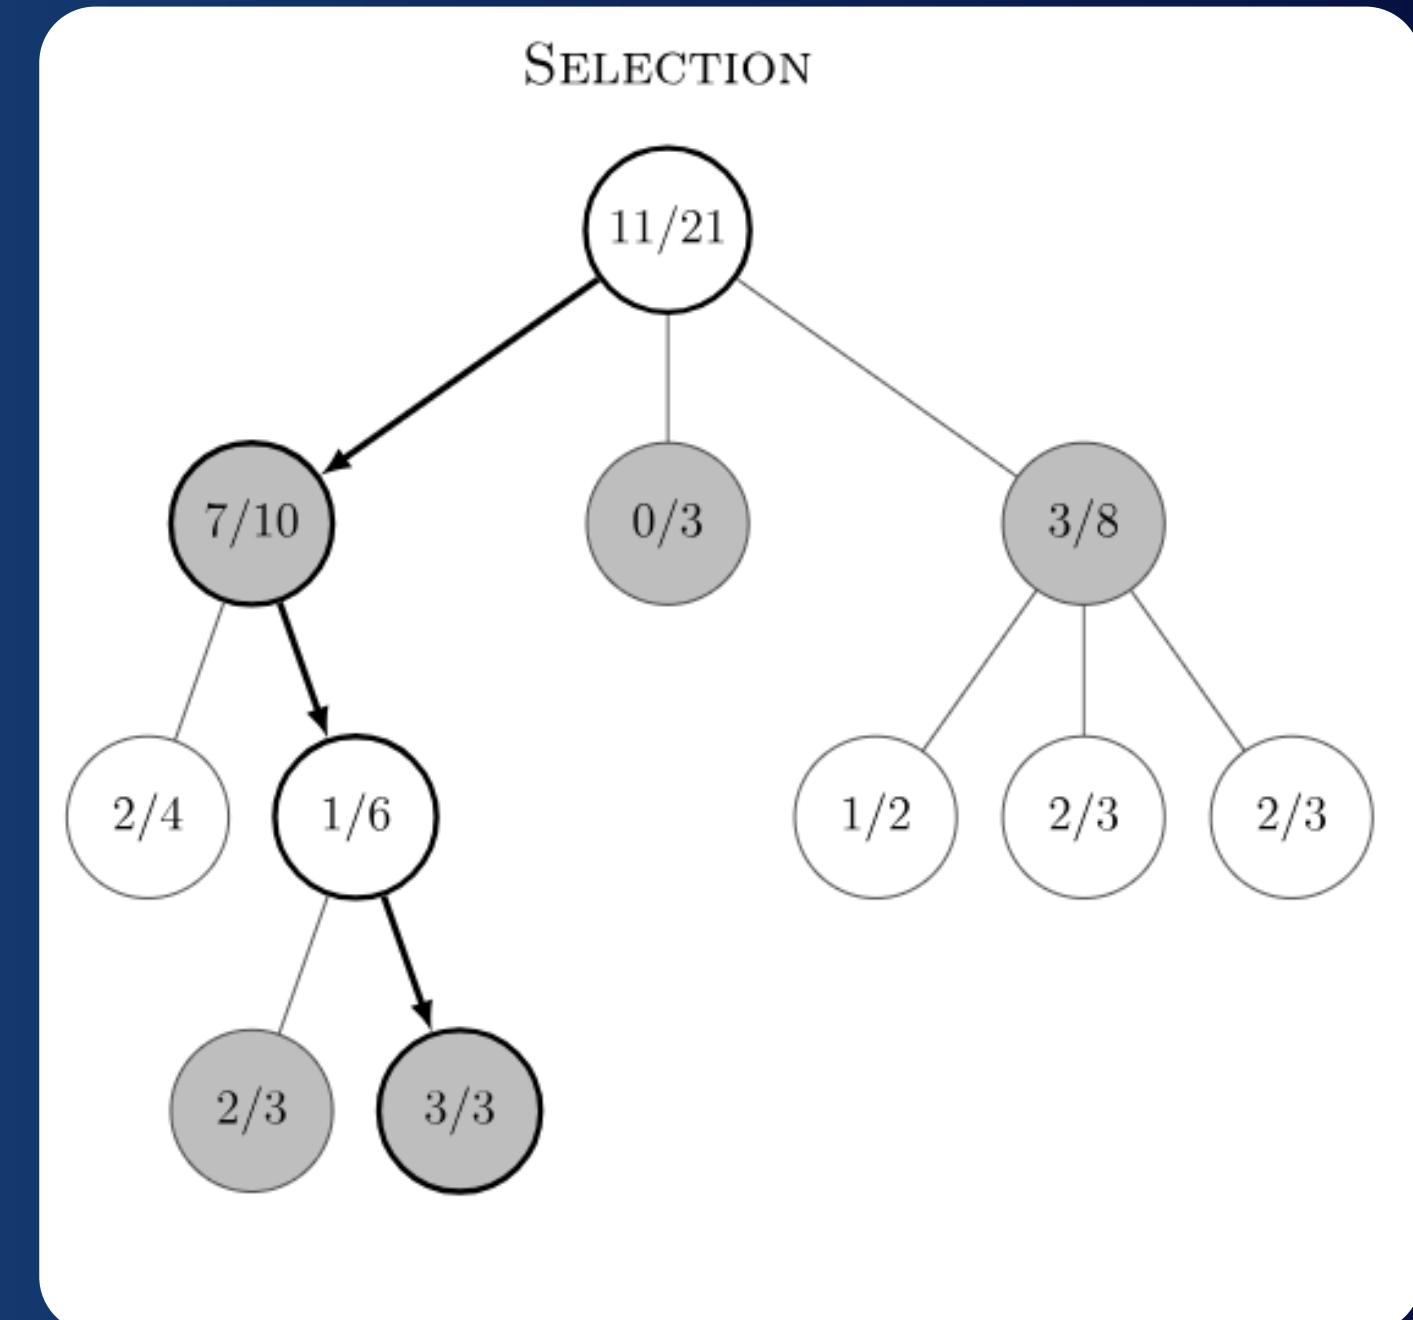


## Backpropagation



# 1- Selecting

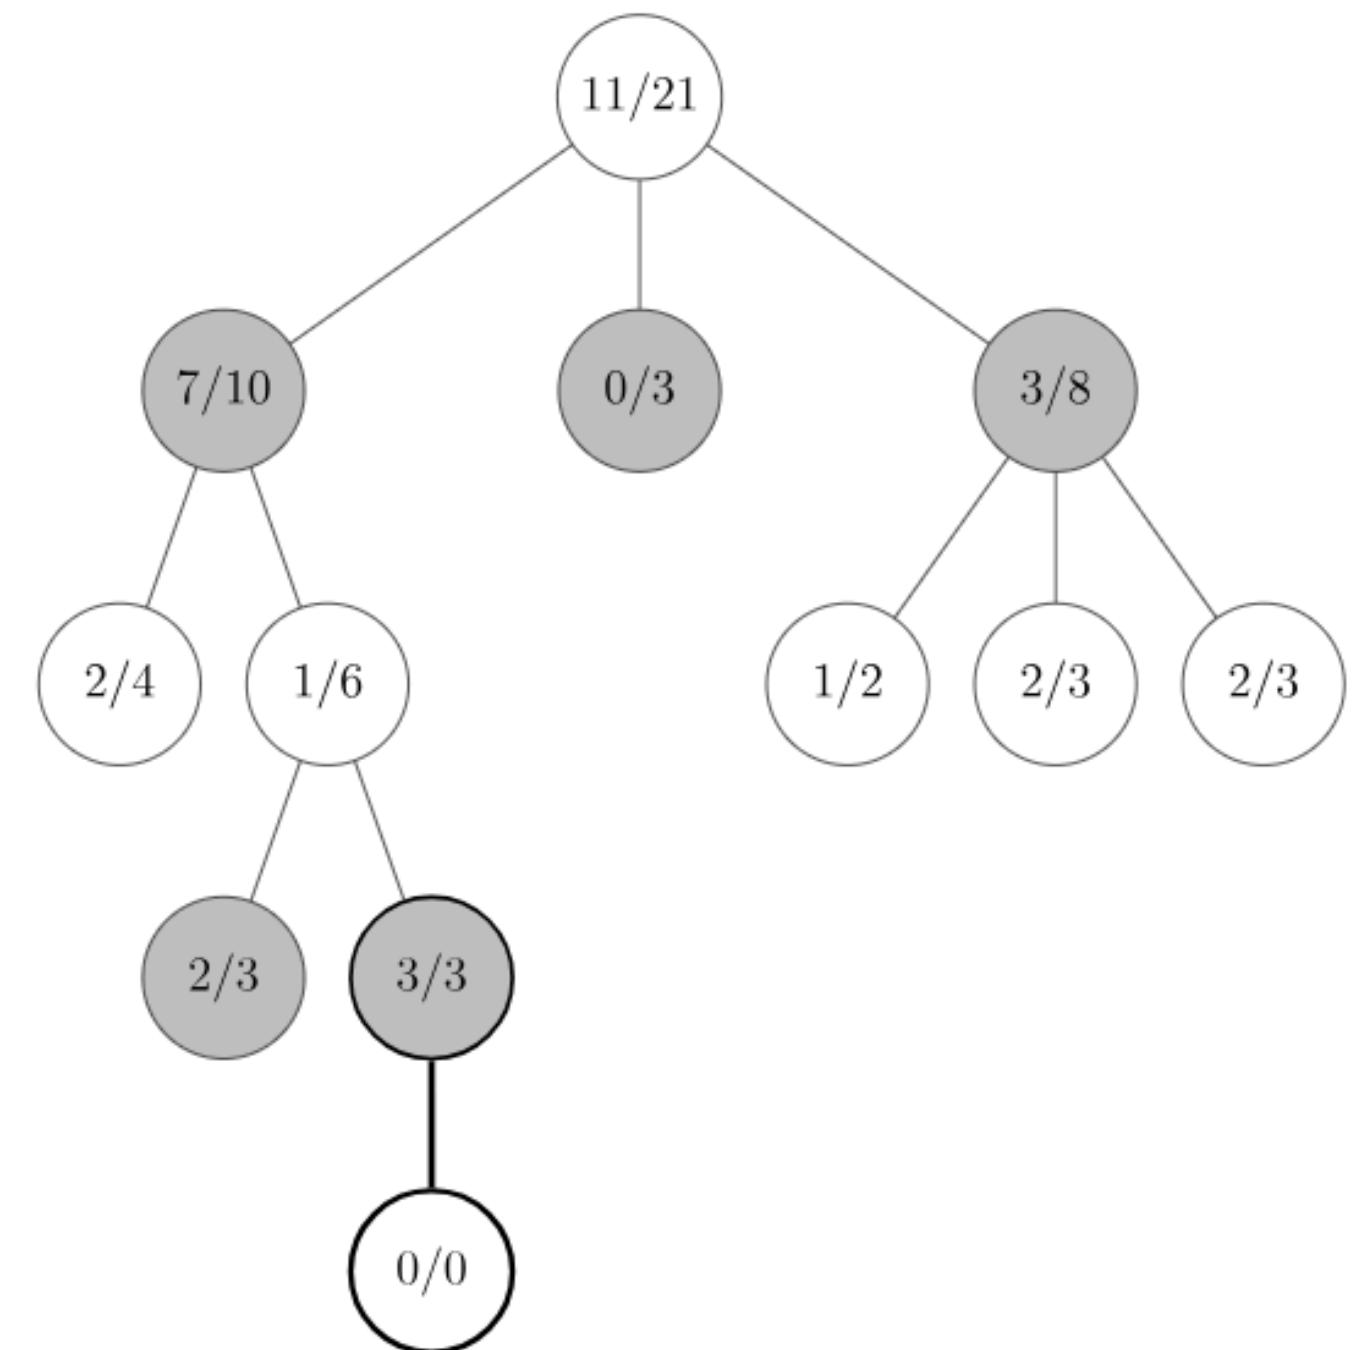
- Start from root R and select successive child nodes until a leaf node L is reached. The root is the current game state and a leaf is any node that has a potential child from which no simulation (playout) has yet been initiated.
- The section says more about a way of biasing choice of child nodes that lets the game tree expand towards the most promising moves, which is the essence of Monte Carlo tree search.



# 2- Expansion

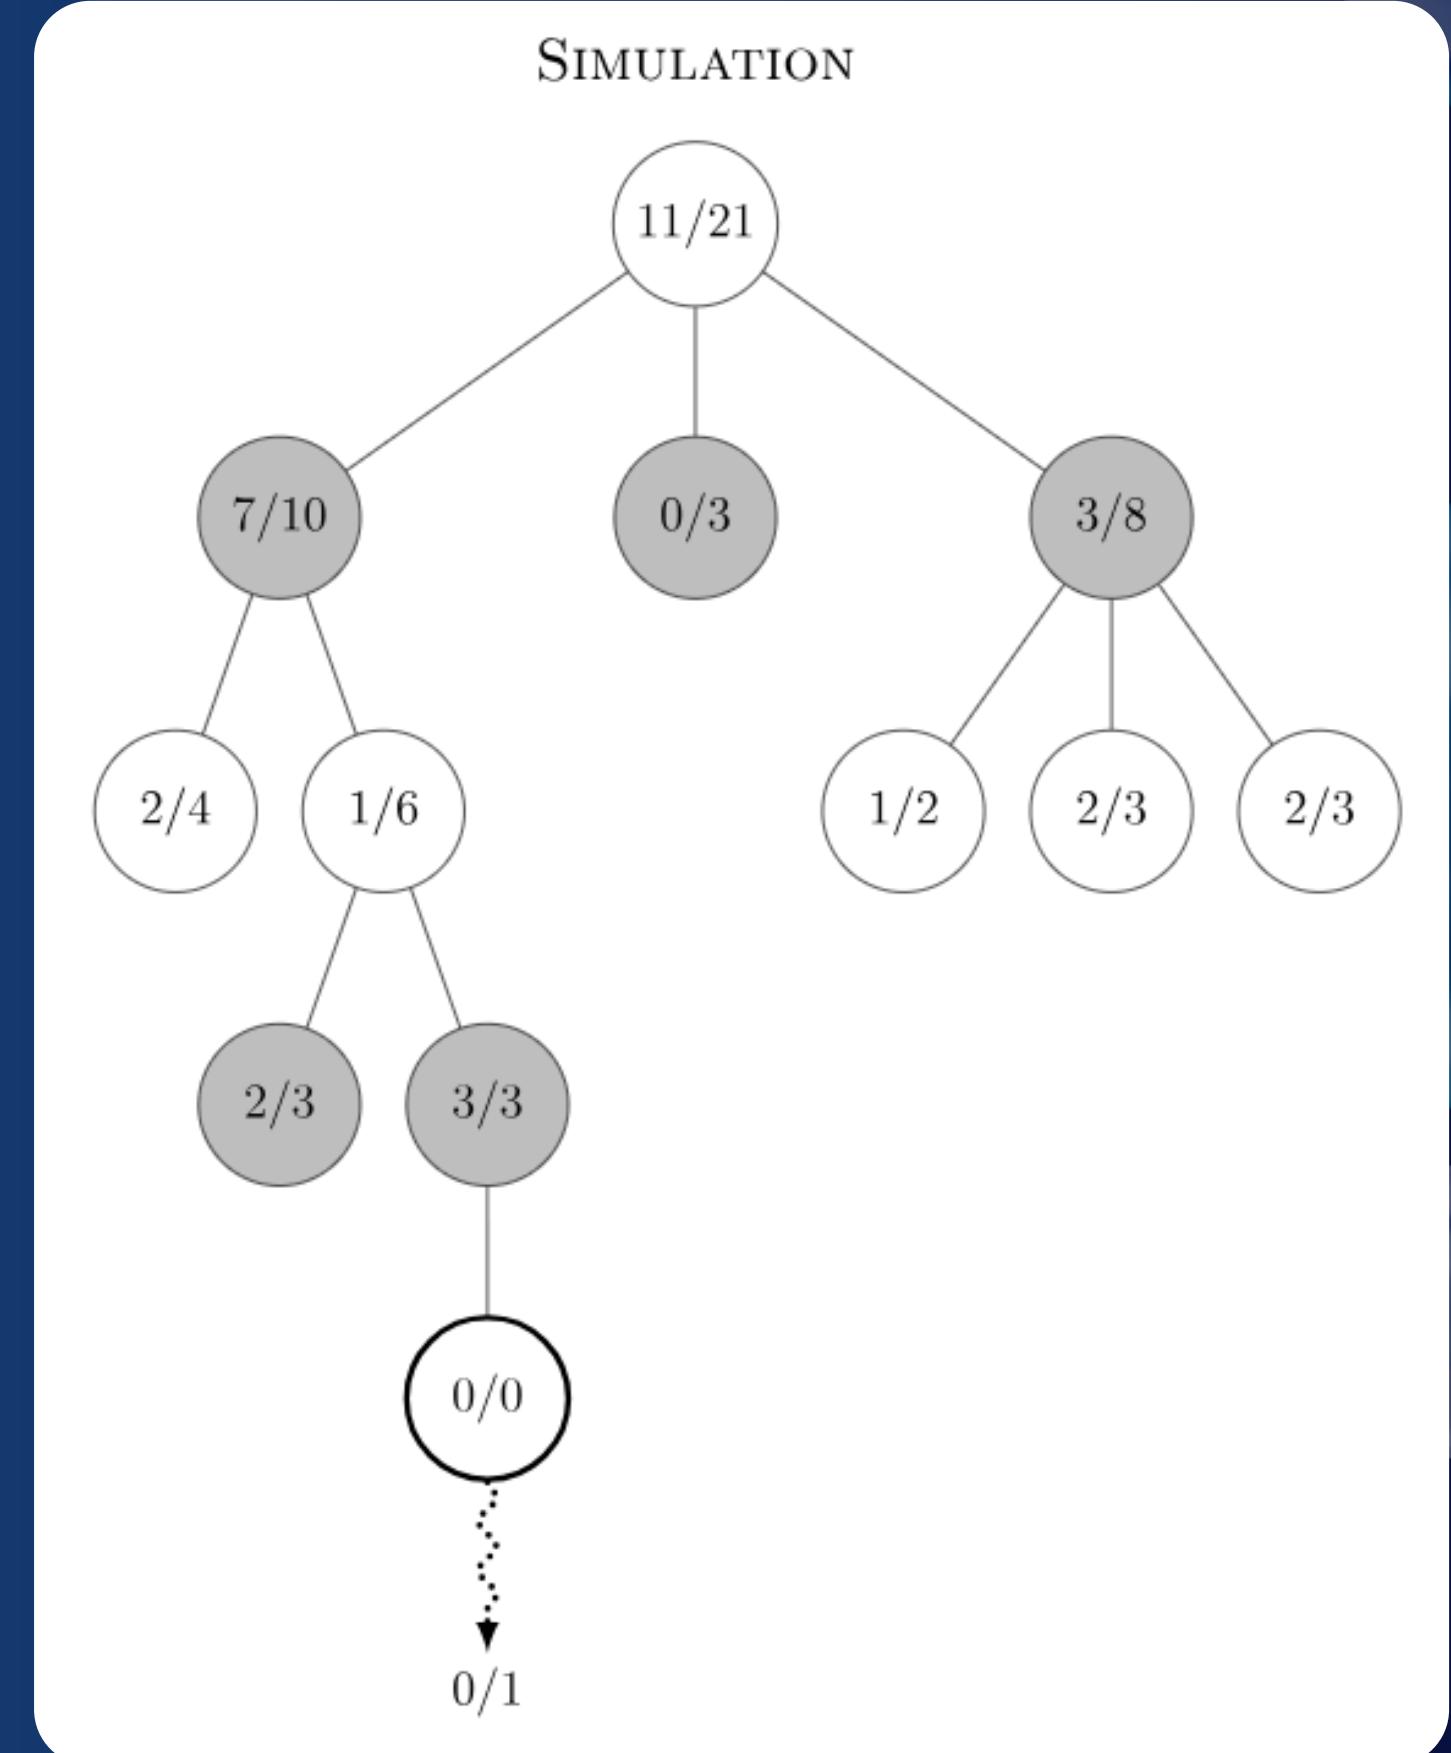
- Unless L ends the game decisively (e.g. win/loss/draw) for either player, create one (or more) child nodes and choose node C from one of them. Child nodes are any valid moves from the game position defined by L.

EXPANSION



# 3- Simulation

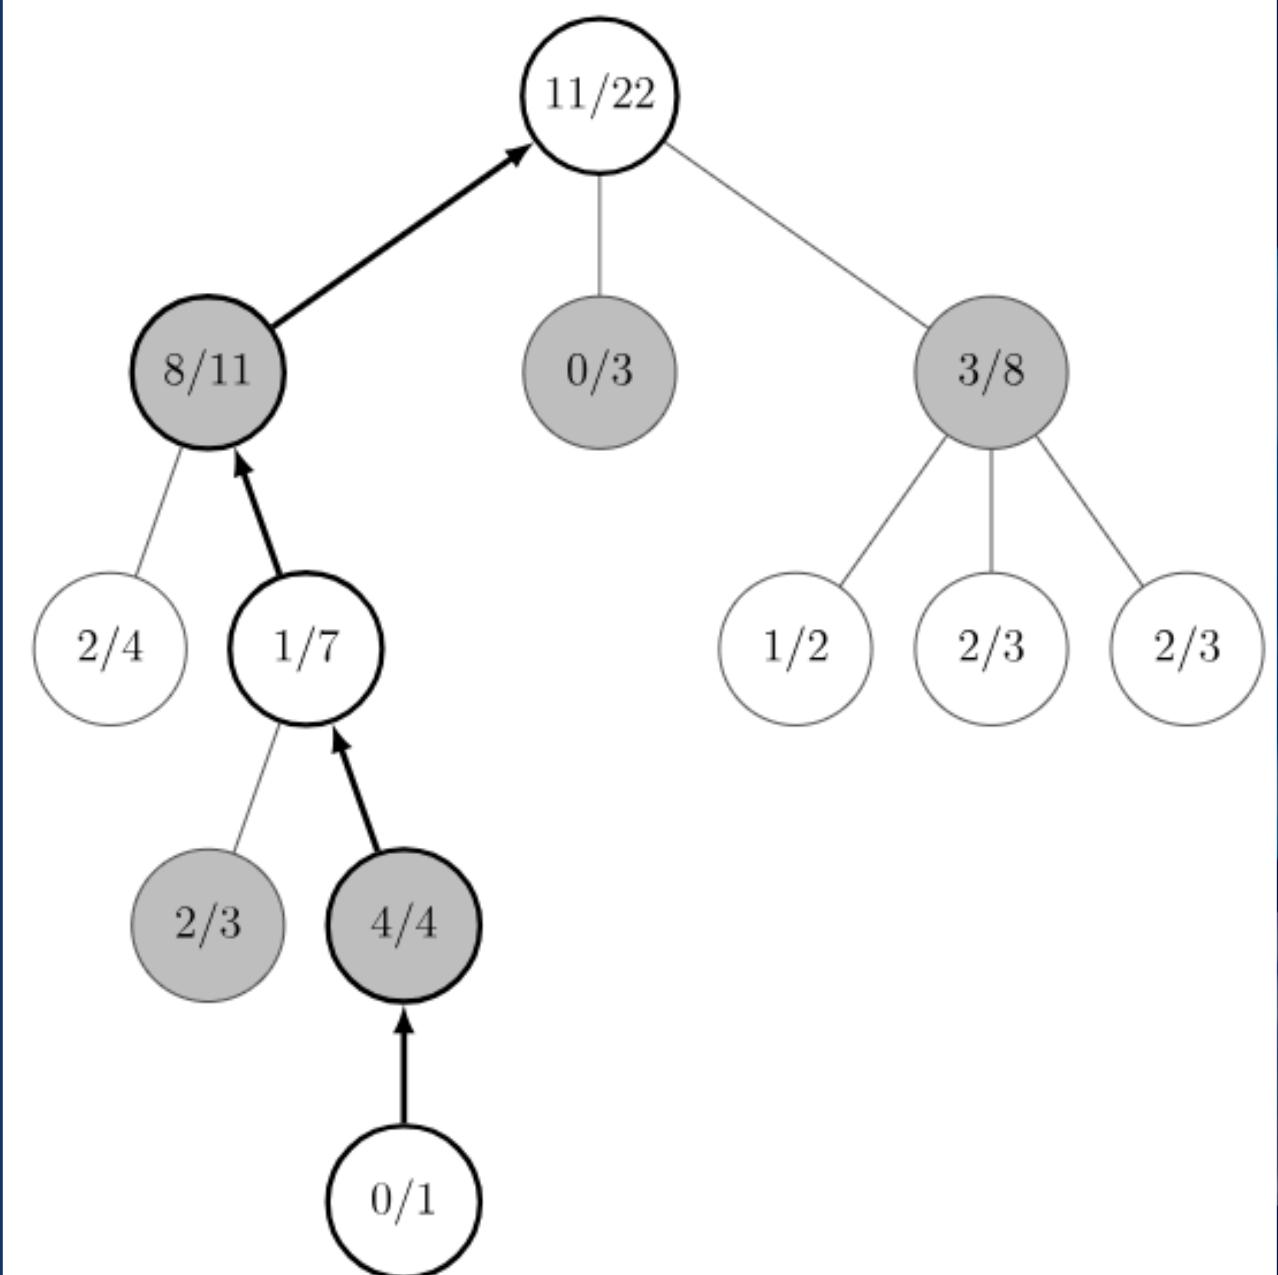
- Complete one random playout from node C. This step is sometimes also called playout or rollout. A playout may be as simple as choosing uniform random moves until the game is decided (for example in chess, the game is won, lost, or drawn).



# 4- Backpropagation

- Use the result of the playout to update information in the nodes on the path from C to R.

BACKPROPAGATION



# Exploration and Exploitation

- The main difficulty in selecting child nodes is maintaining some **balance between the exploitation** of deep variants after moves with high average win rate and **the exploration of moves** with few simulations.

- The first formula for balancing exploitation and exploration in games, called **UCT (Upper Confidence Bound )**

Monte Carlo tree search



# UCT Formula (Upper Confidence Bound)

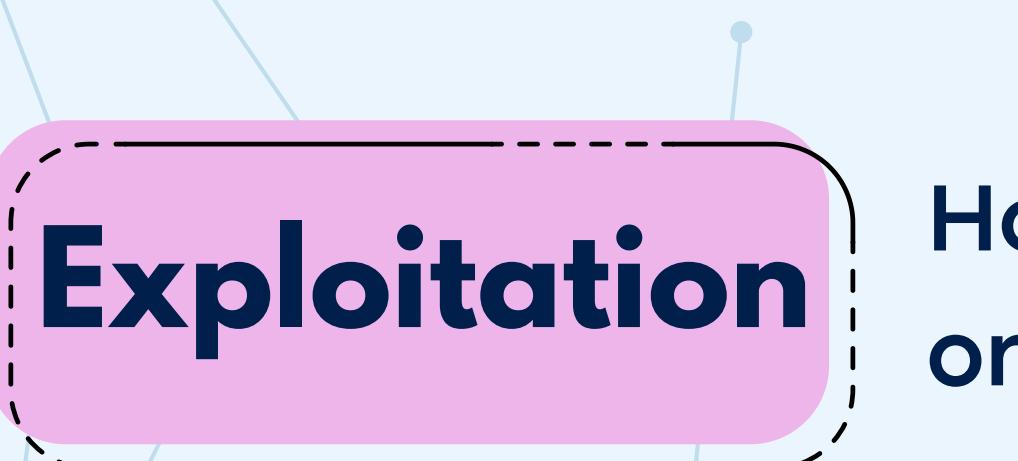
$$UCT = \frac{w_i}{n_i} + c \sqrt{\frac{\ln N_i}{n_i}}$$

- $w_i$  → stands for the number of wins for the node considered after the  $i$ -th move
- $n_i$  → stands for the number of simulations for the node considered after the  $i$ -th move
- $N_i$  → stands for the total number of simulations after the  $i$ -th move run by the parent node of the one considered
- $c$  → is the exploration parameter—theoretically equal to  $\sqrt{2}$ ; in practice usually chosen empirically

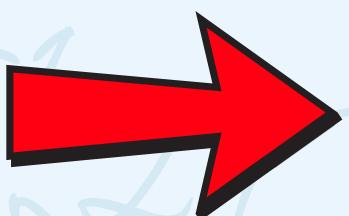
# UCT Formula (Upper Confidence Bound)



The UCT formula decides which child node to explore next by combining:



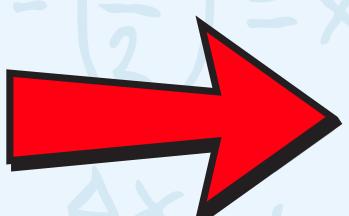
How good is this move based  
on past simulations?



$$\frac{w_i}{n_i}$$



How much we haven't tried  
this move yet?



$$c \cdot \sqrt{\frac{\ln N}{n_i}}$$

Monte Carlo tree search



## Balanced Decision Making:

- If a node has been visited a lot, its exploration term **becomes smaller** → it won't be favored unless it has high win rate.
- If a node has been visited rarely, the exploration term **becomes larger** → we encourage trying new paths.



## Why is it important?

Without UCT, MCTS might either:

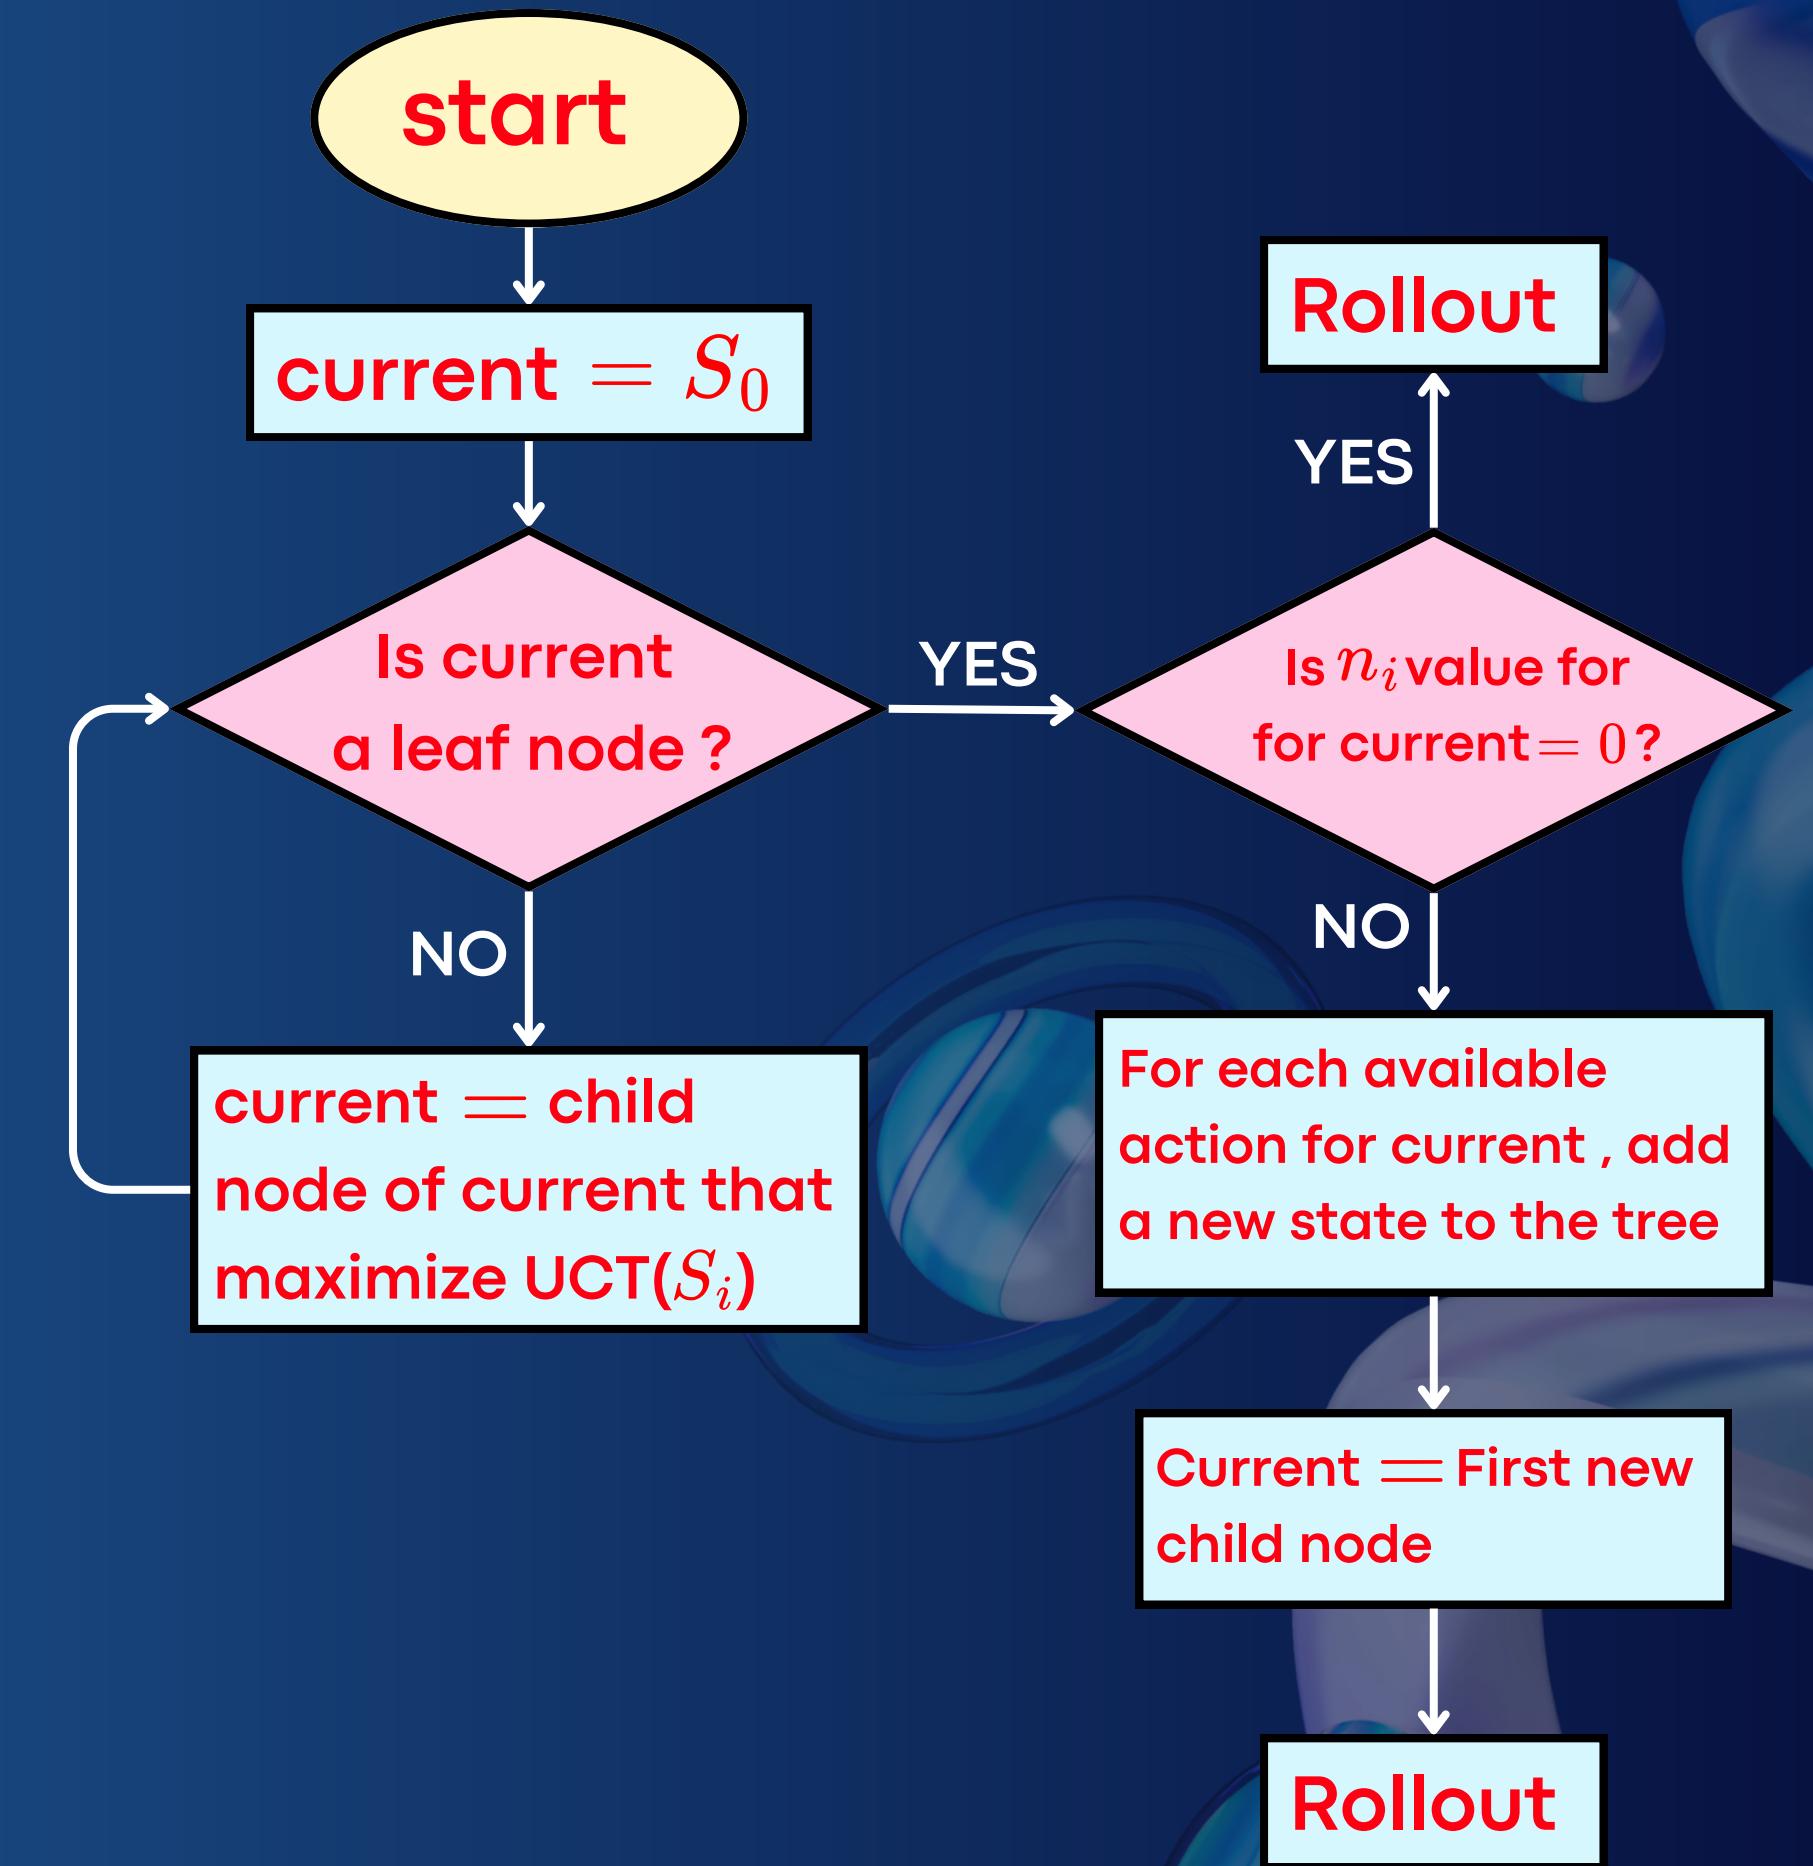
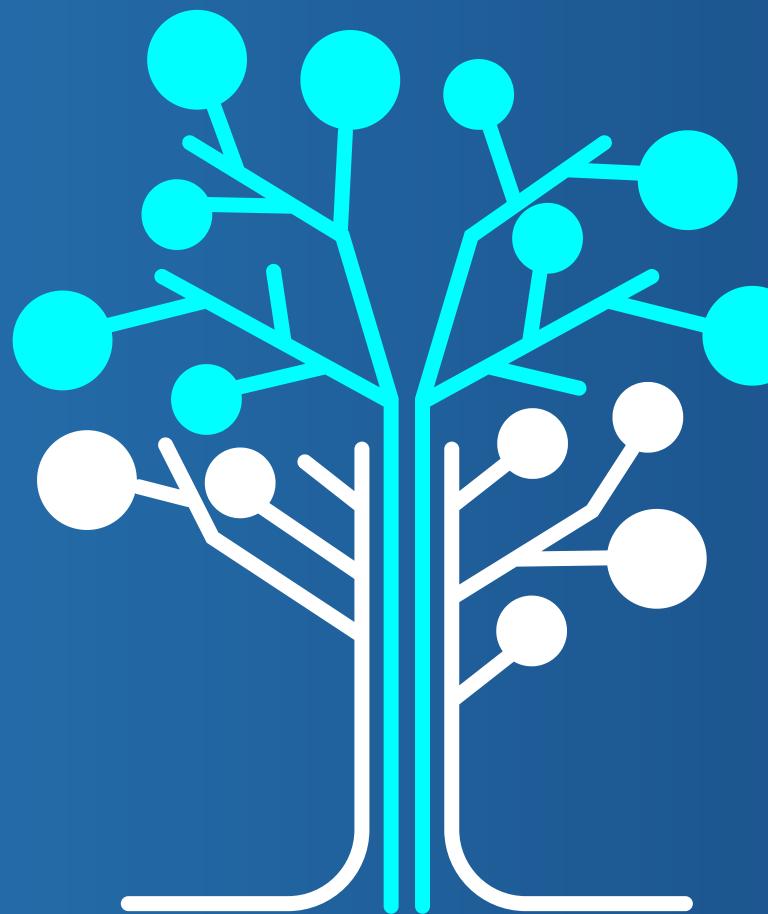
- Get stuck exploiting a move that looked good early on (without trying others).
- Or waste time exploring too many bad moves.

UCT gives the best of both worlds. 

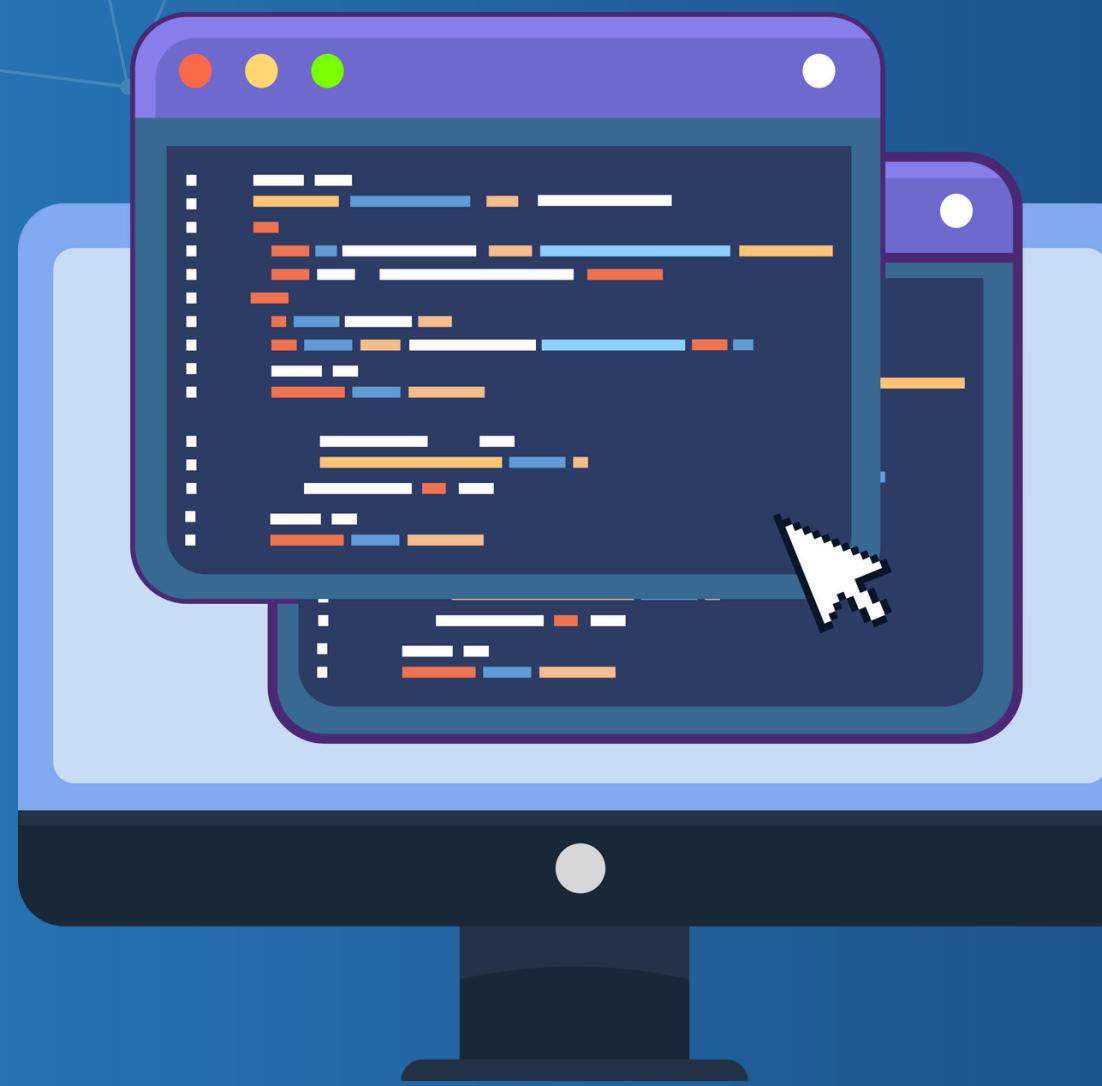


# FLOWCHART

Algorithm for tree traversal  
and node Expansion



# Pseudocode



Terminal  
state

```
function TreeTraversalAndExpansion(root):
    current ← root // S0

    while current is not a leaf node:
        current ← child of current that maximizes UCT

        if n(current) = 0: // node has never been visited
            return Rollout(current)

        else:
            for each available action from current:
                new_state ← result of applying action
                Add new_state as a child node to current

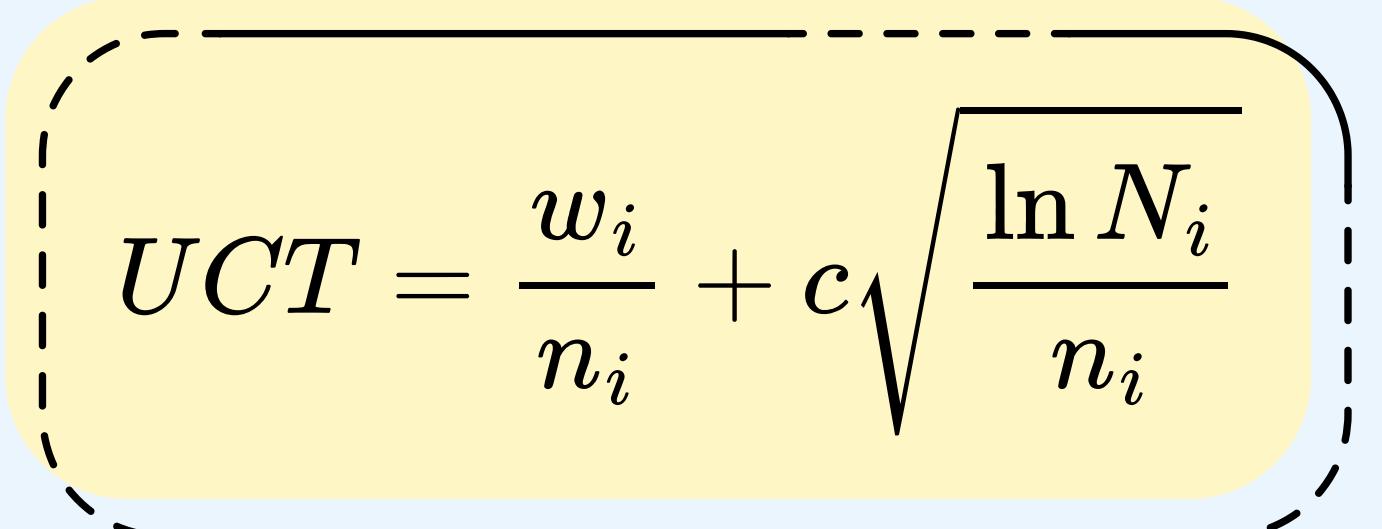
            current ← first new child node added
            return Rollout(current)

function Rollout(si):
    loop forever:
        if si is a terminal state:
            return value(si)

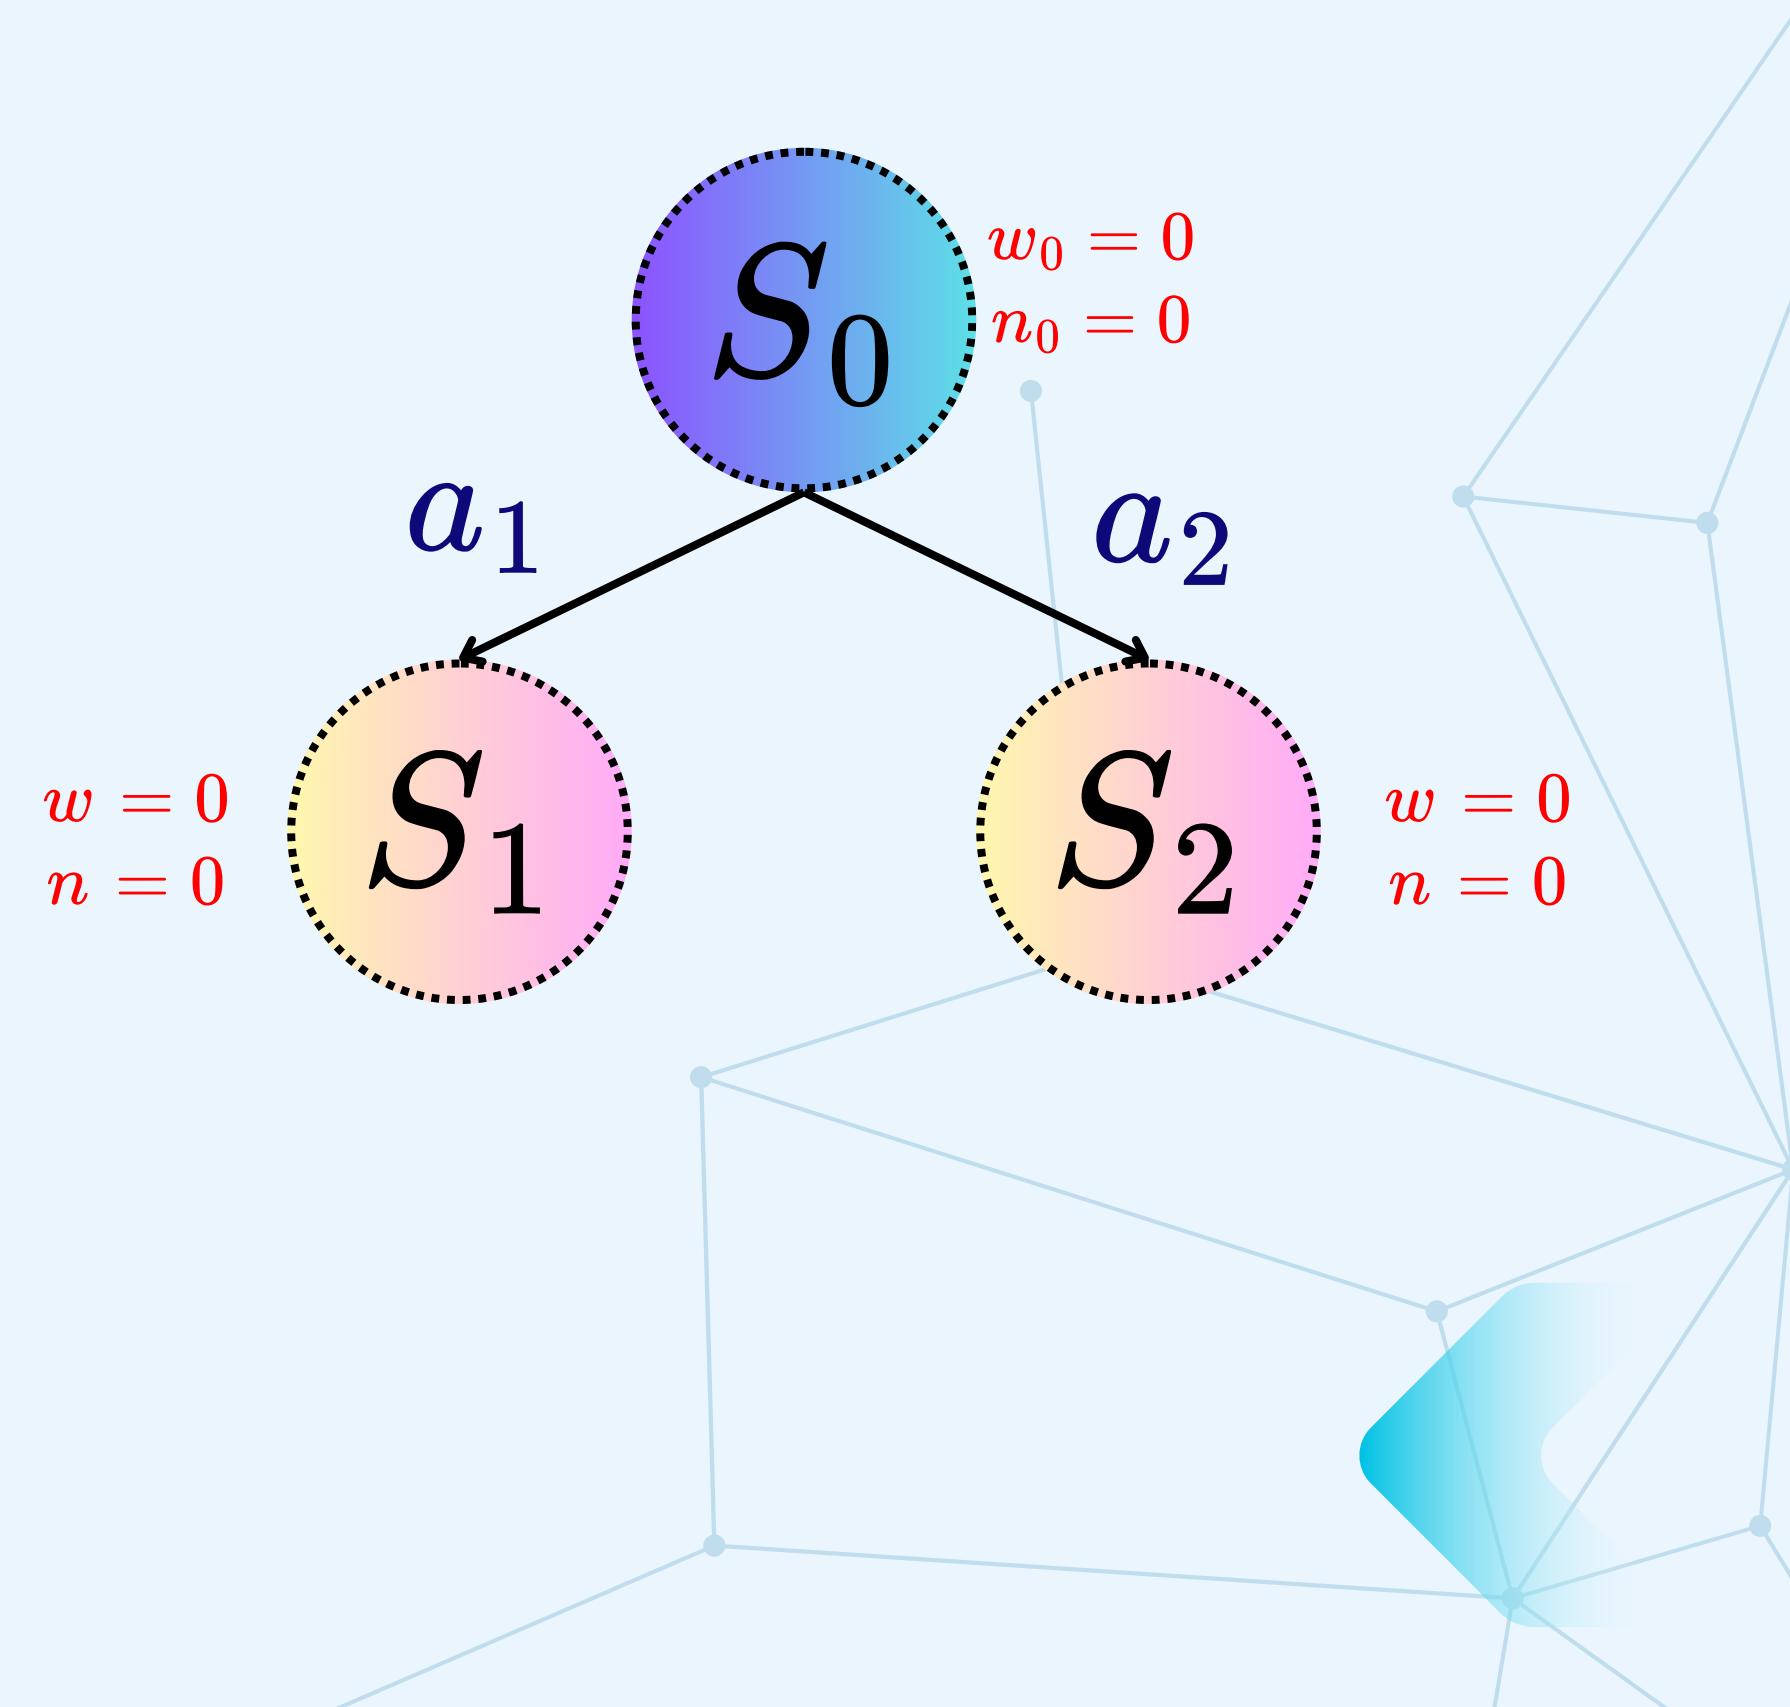
        ai ← random choice from available actions of si
        si ← simulate(si, ai)
```



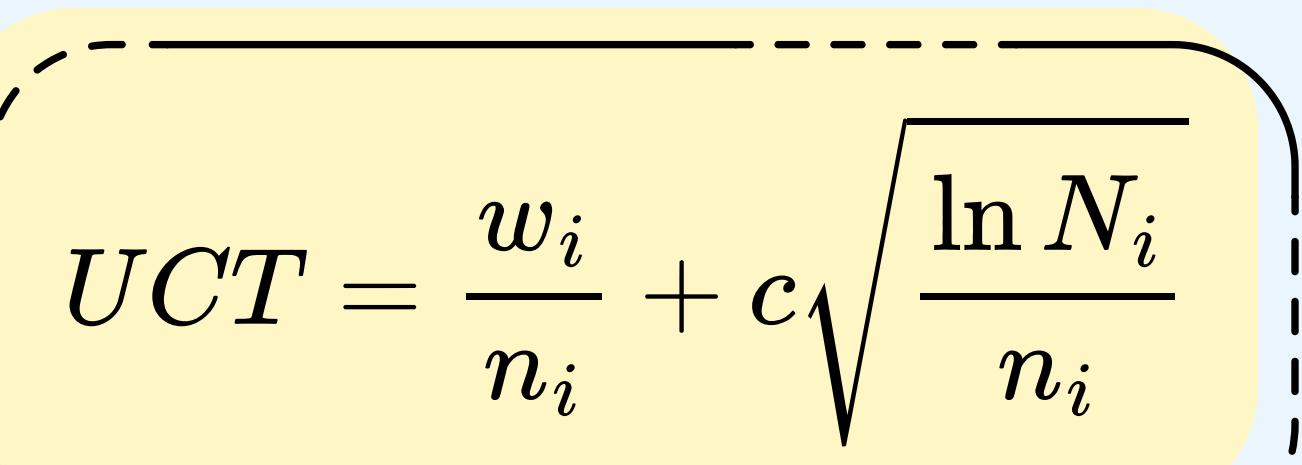
# Initial state

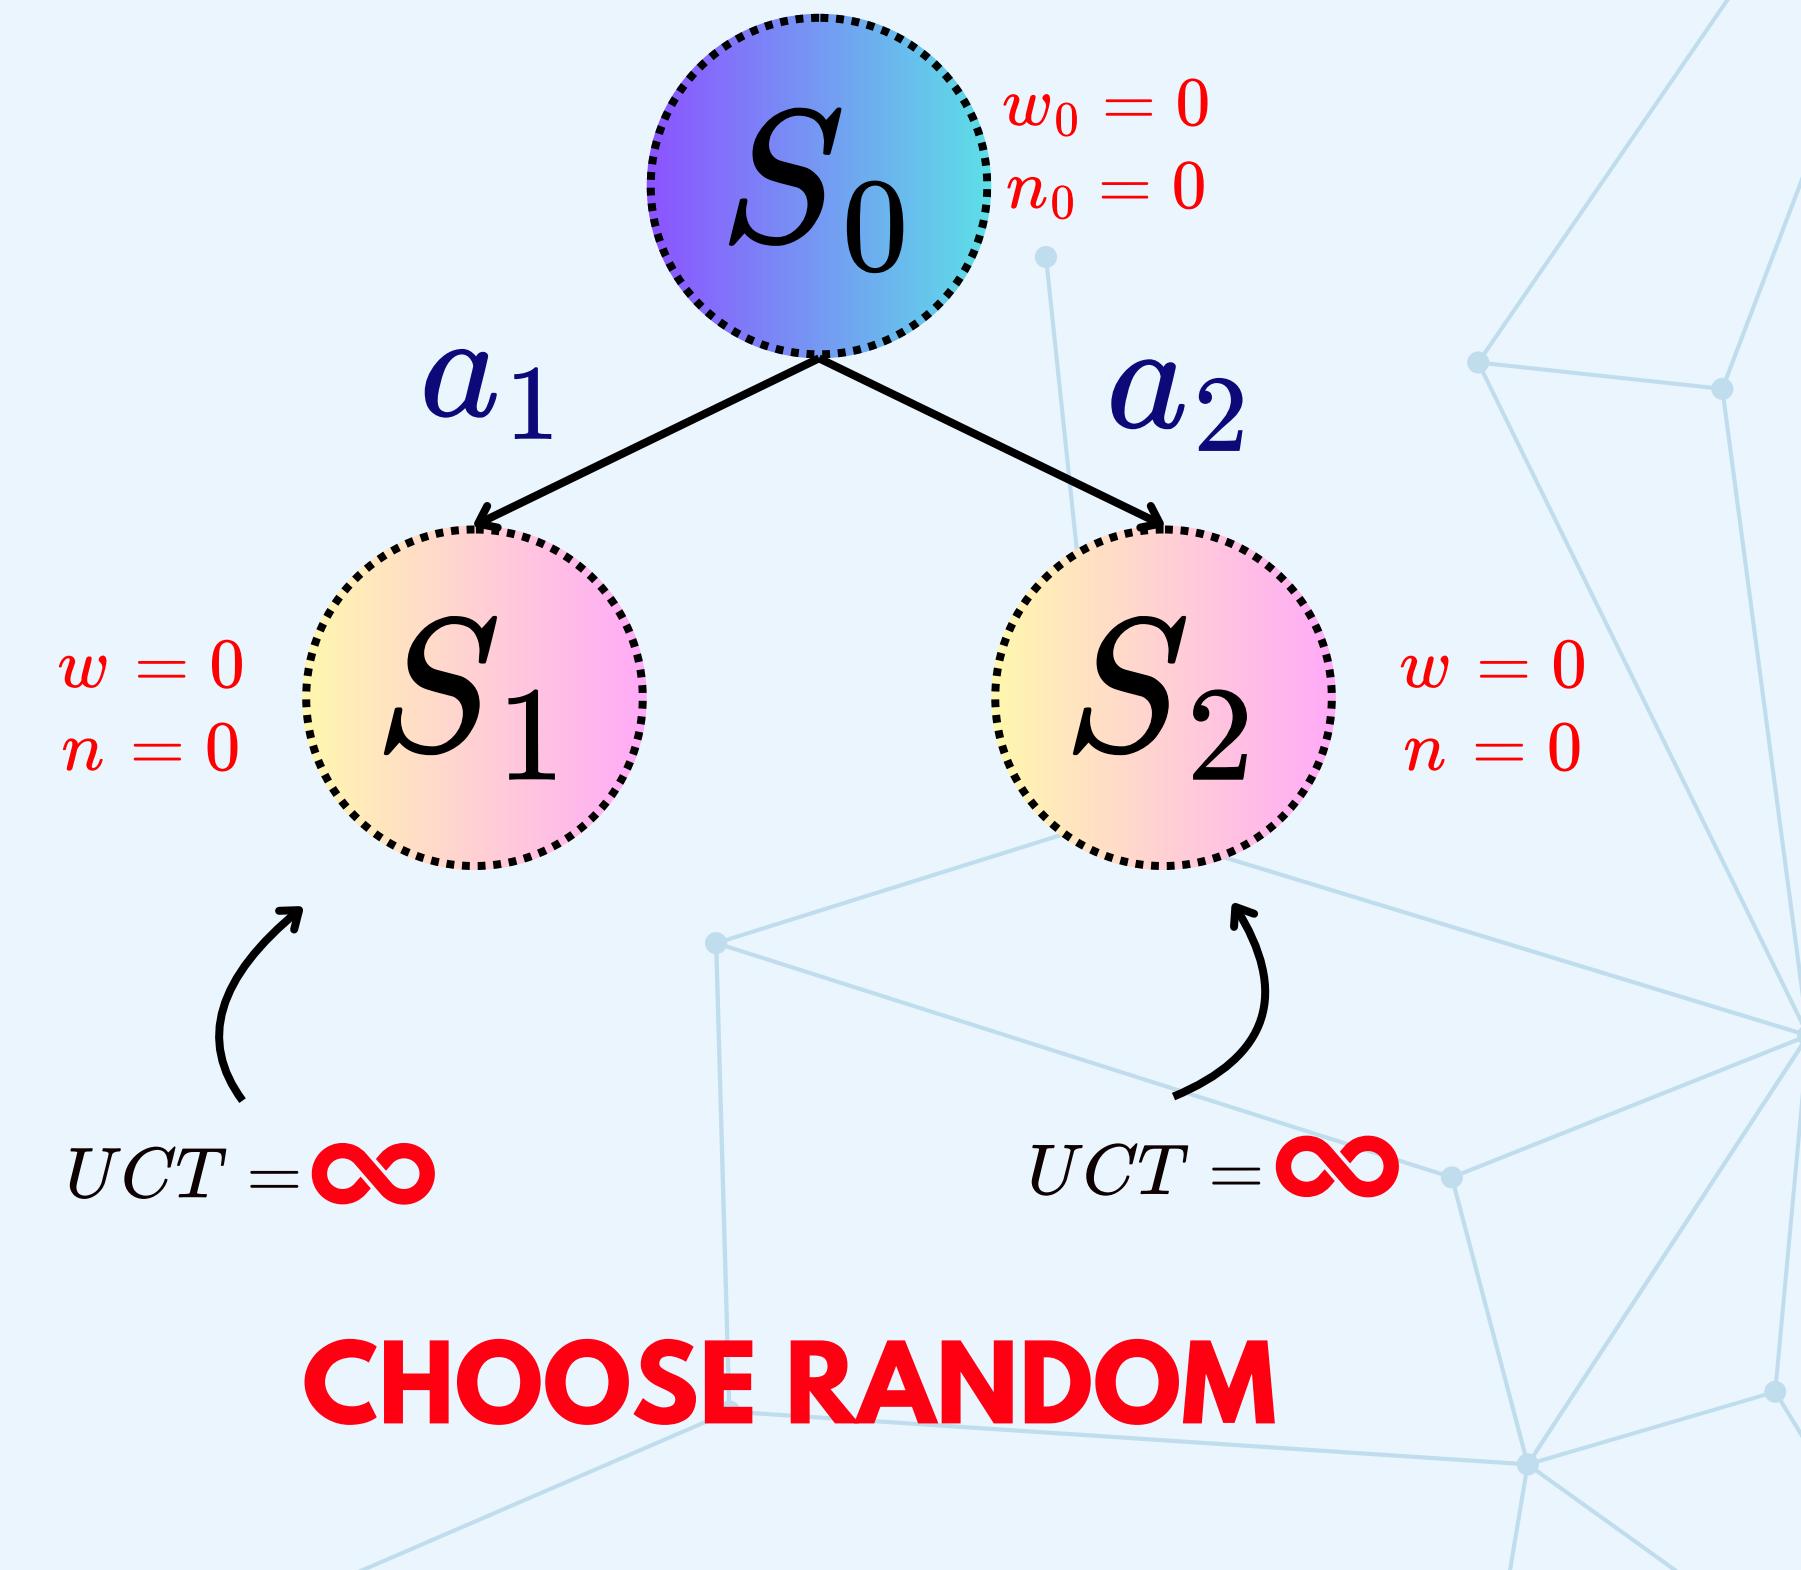

$$UCT = \frac{w_i}{n_i} + c \sqrt{\frac{\ln N_i}{n_i}}$$


Monte Carlo tree search

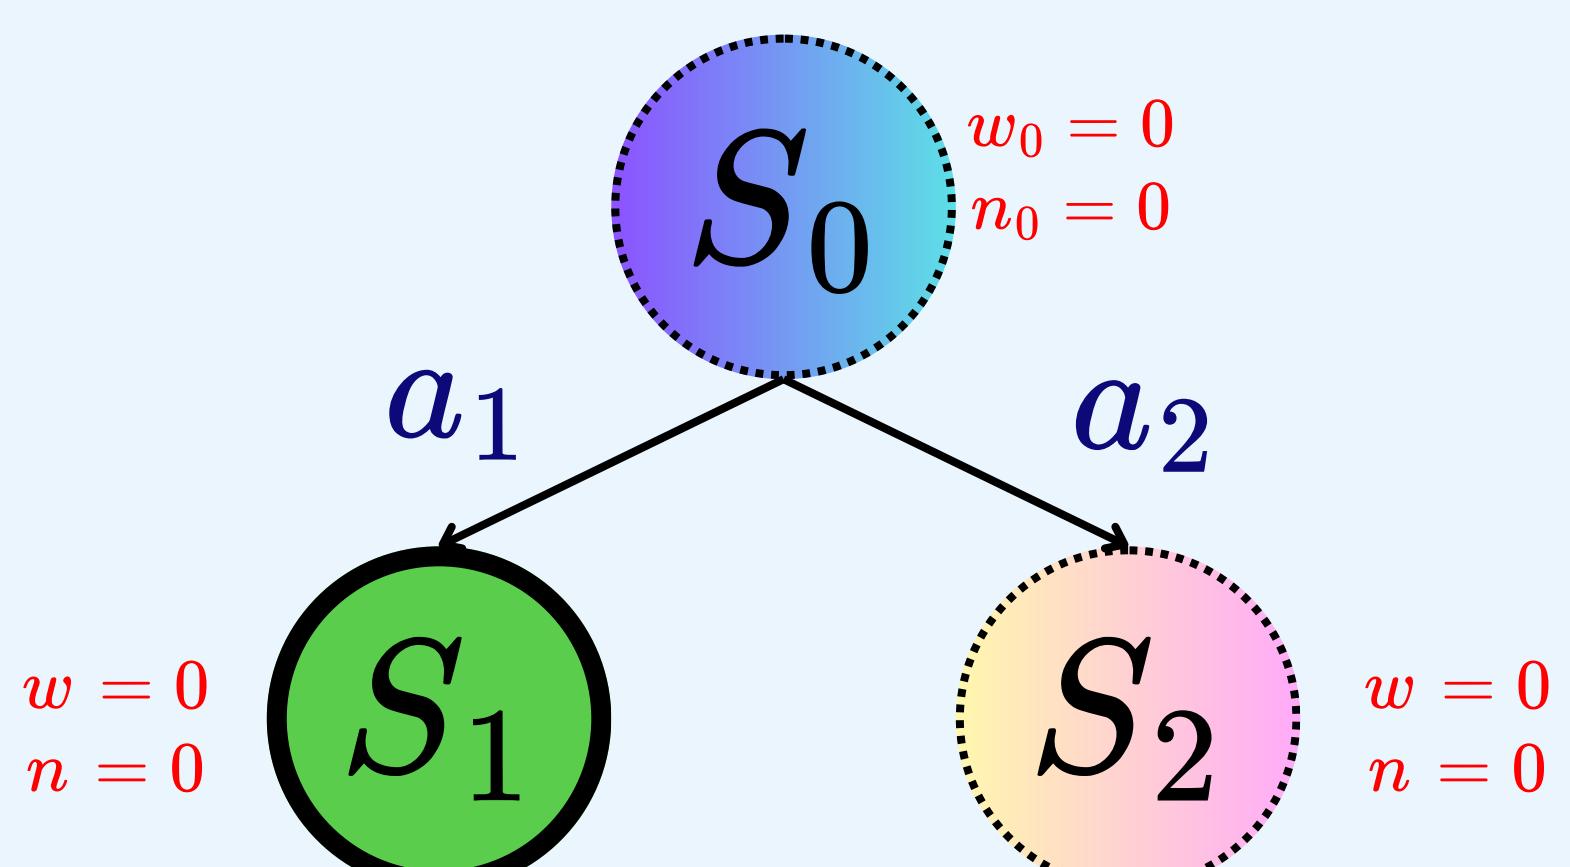


# First Iteration

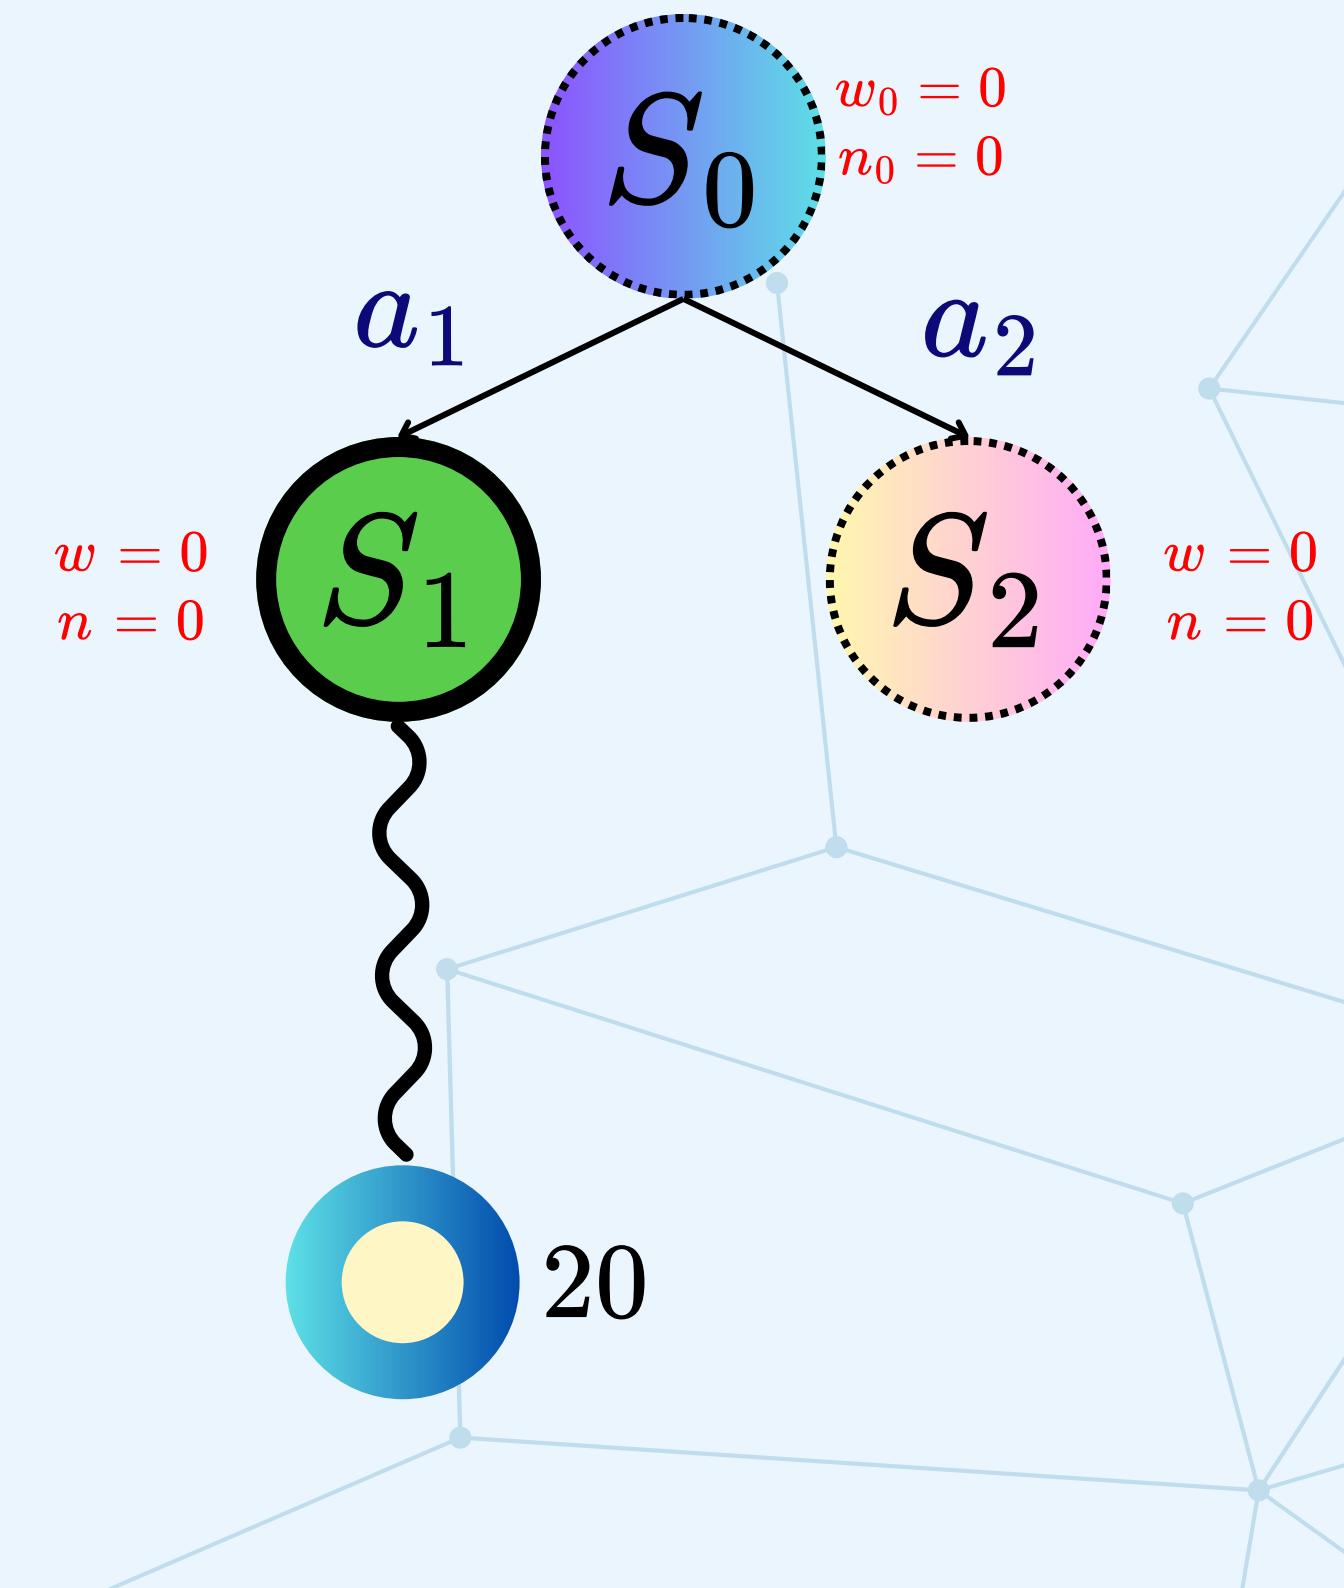

$$UCT = \frac{w_i}{n_i} + c \sqrt{\frac{\ln N_i}{n_i}}$$



# First Iteration

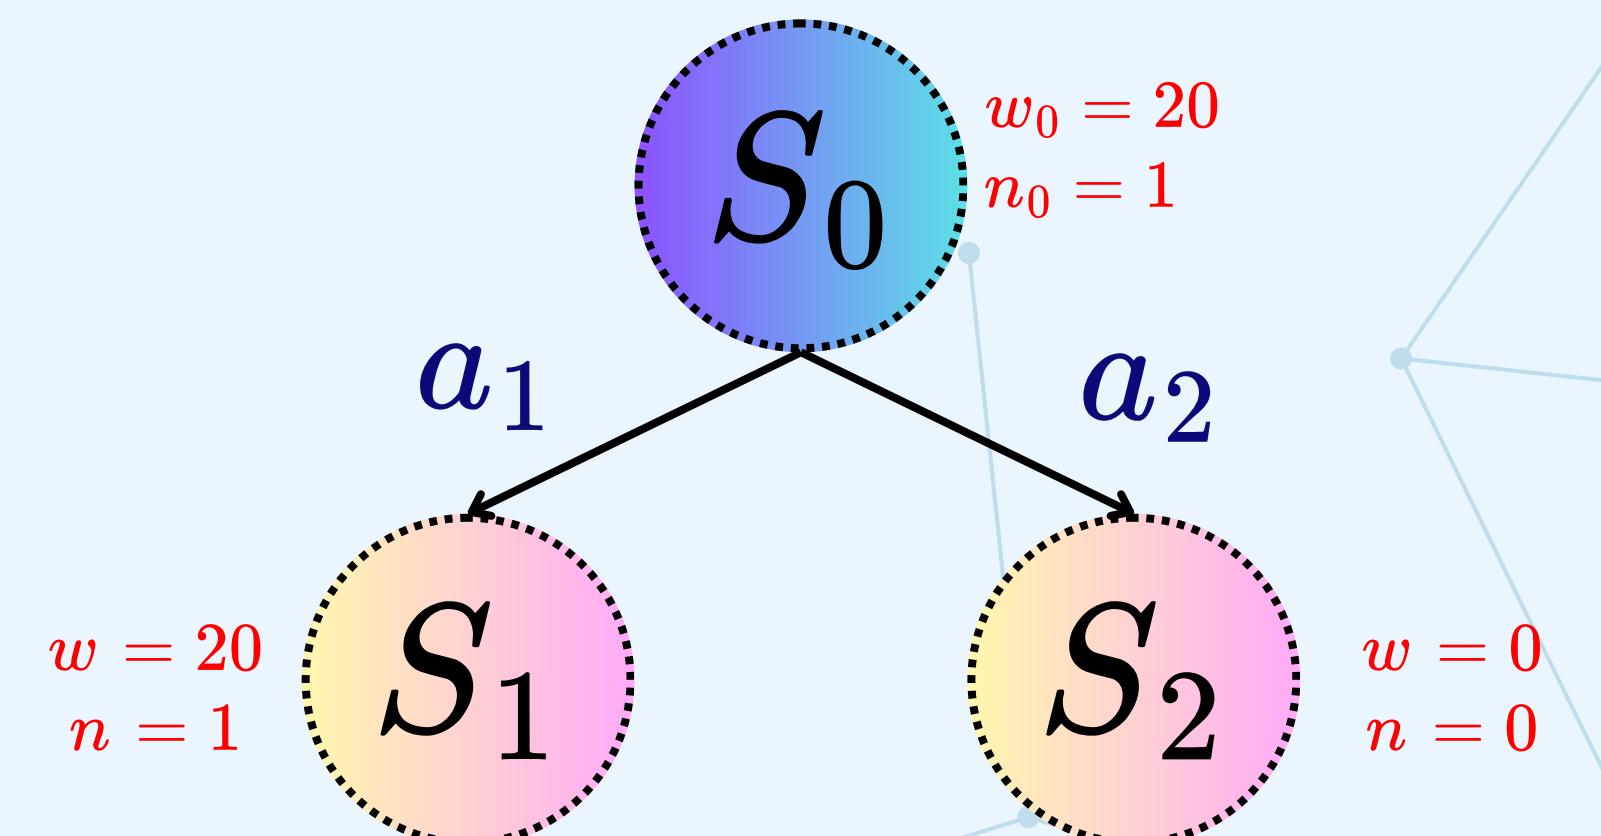
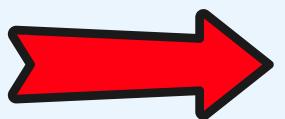
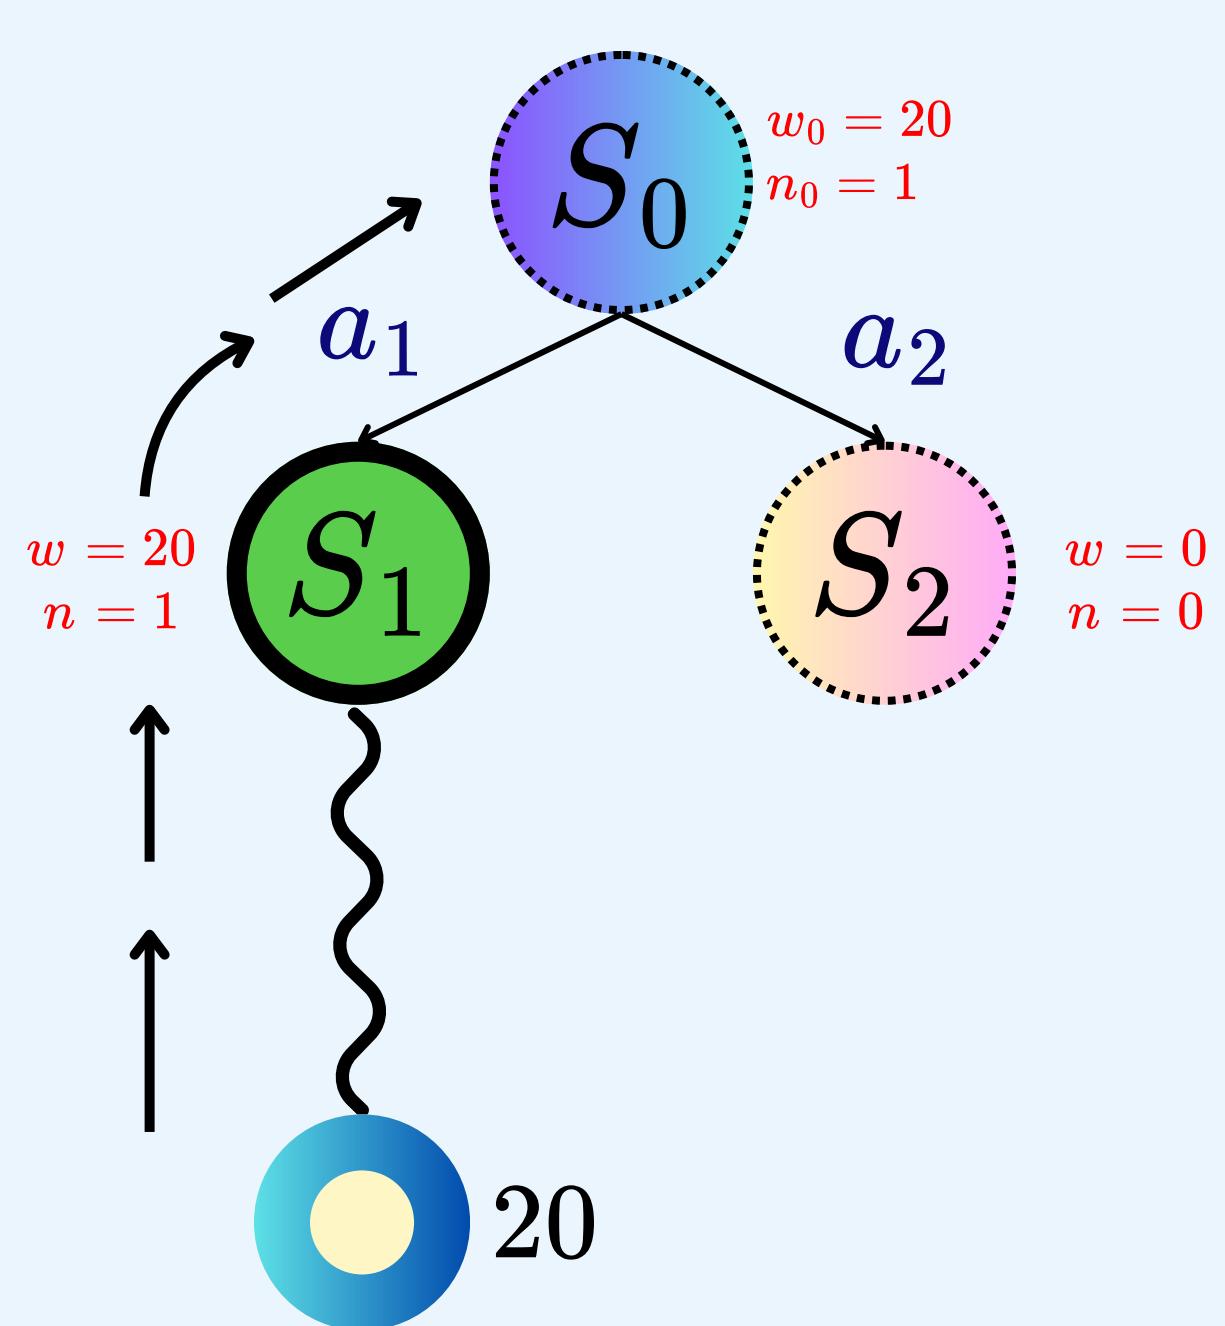


**ROLLOUT !**



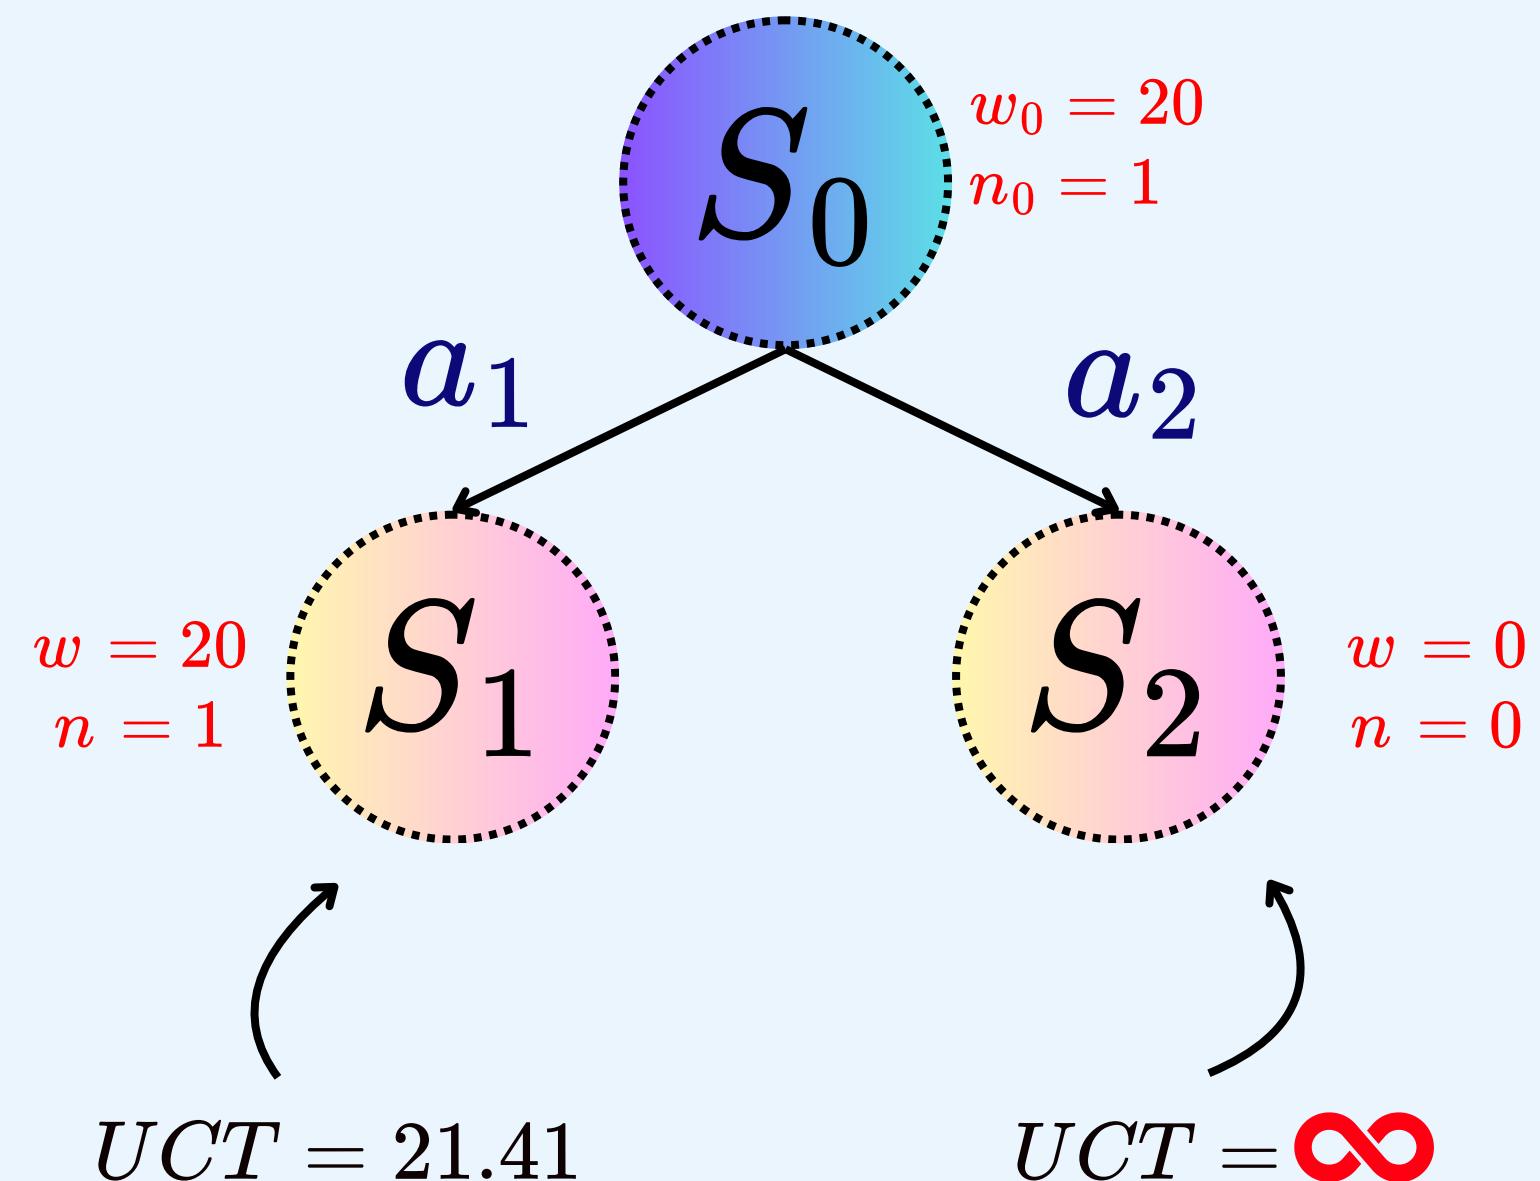
Since  $S_1$  has not been visited yet

# First Iteration

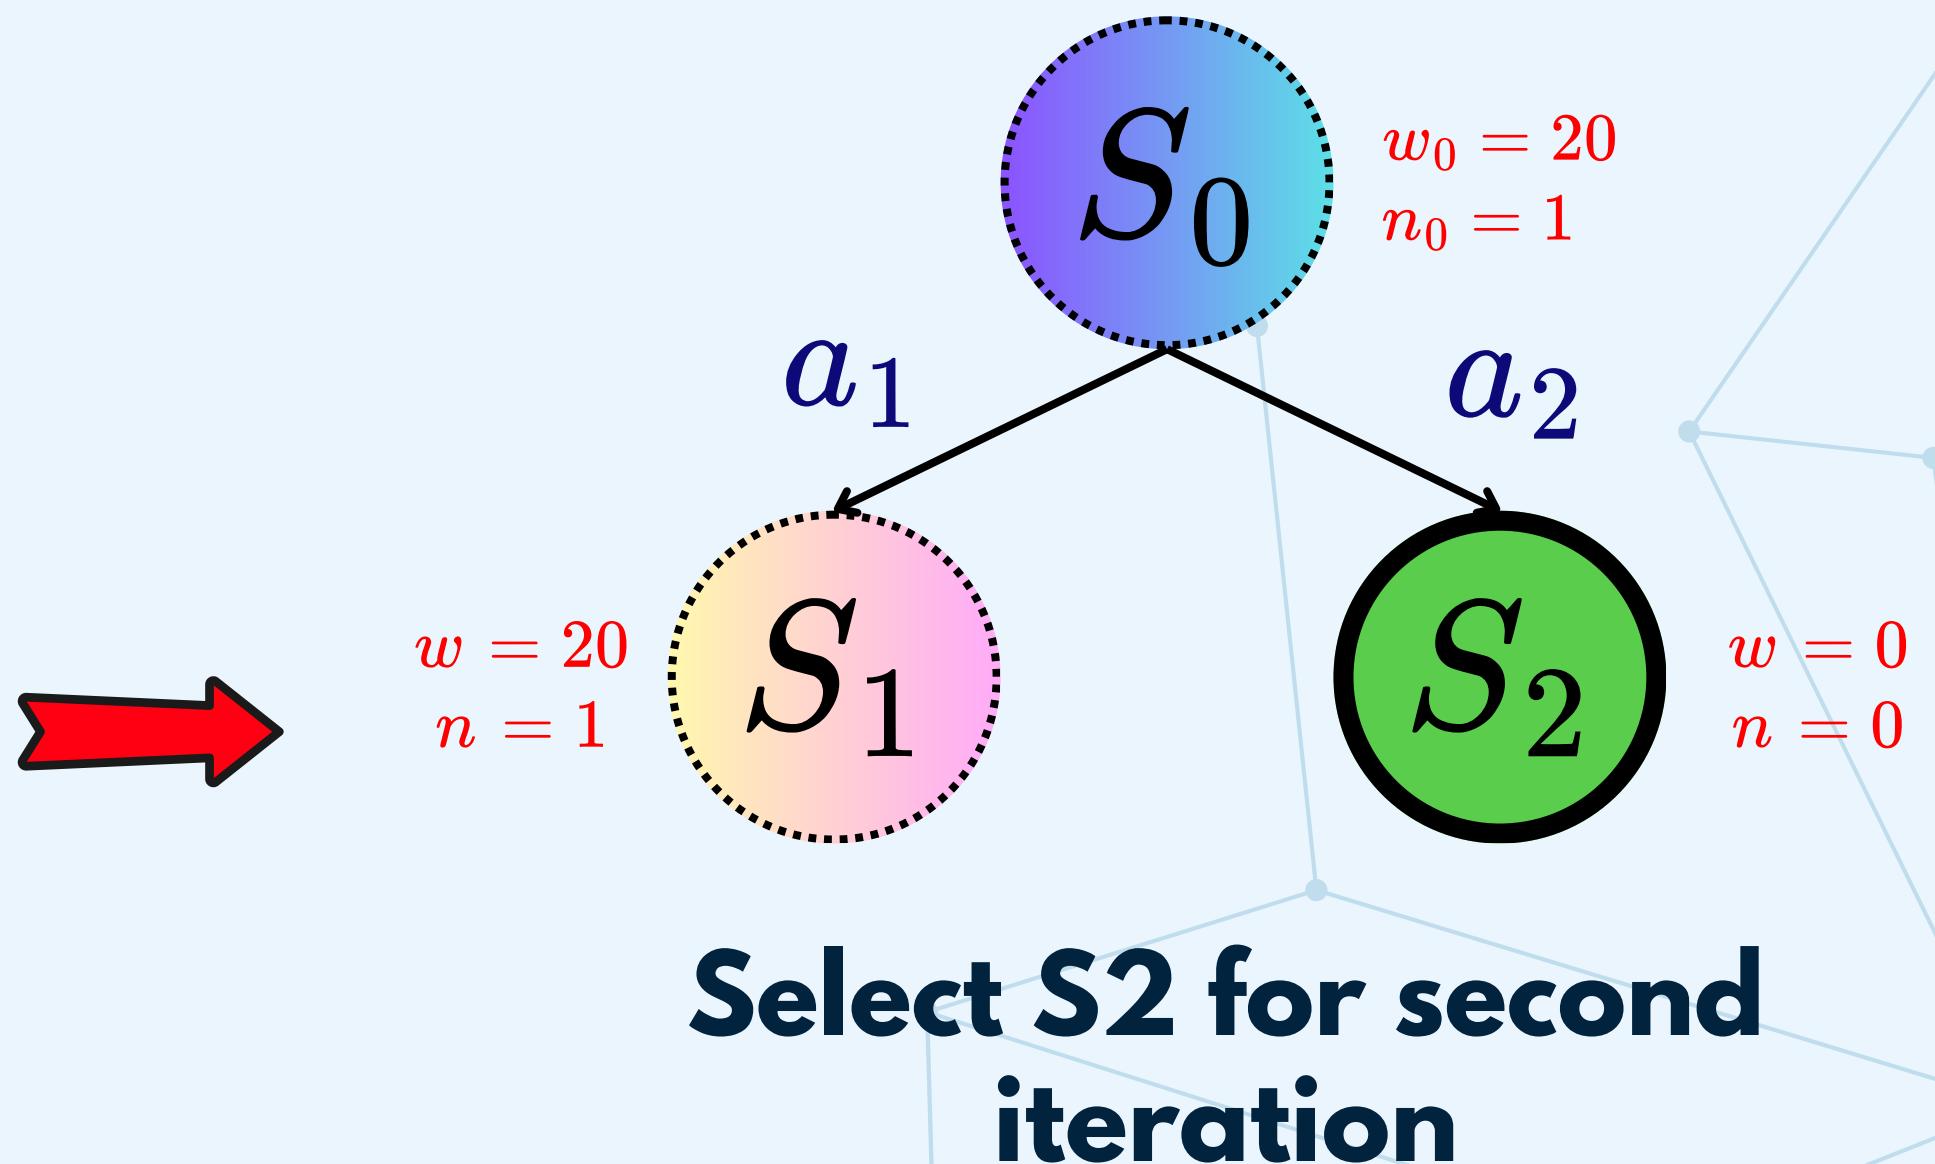


## End of first iteration

## Second Iteration

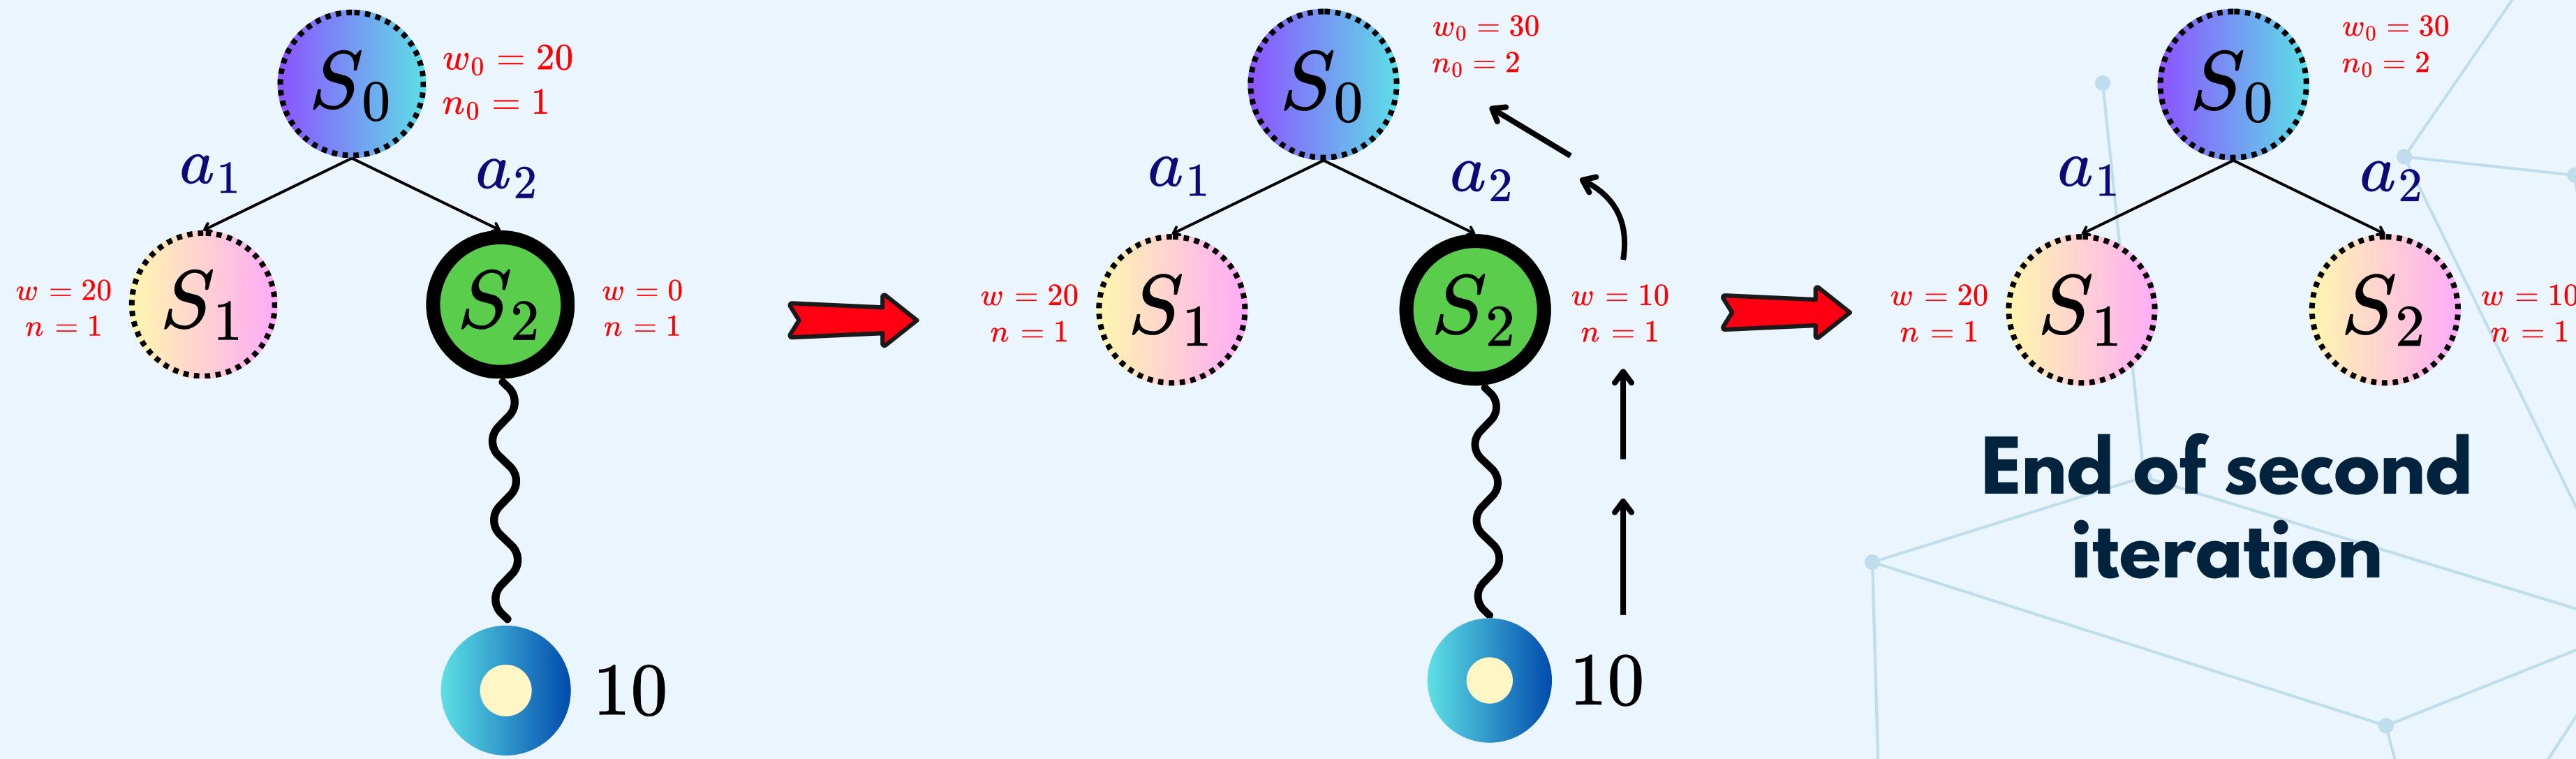


$$UCT = \frac{w_i}{n_i} + c \sqrt{\frac{\ln N_i}{n_i}}$$

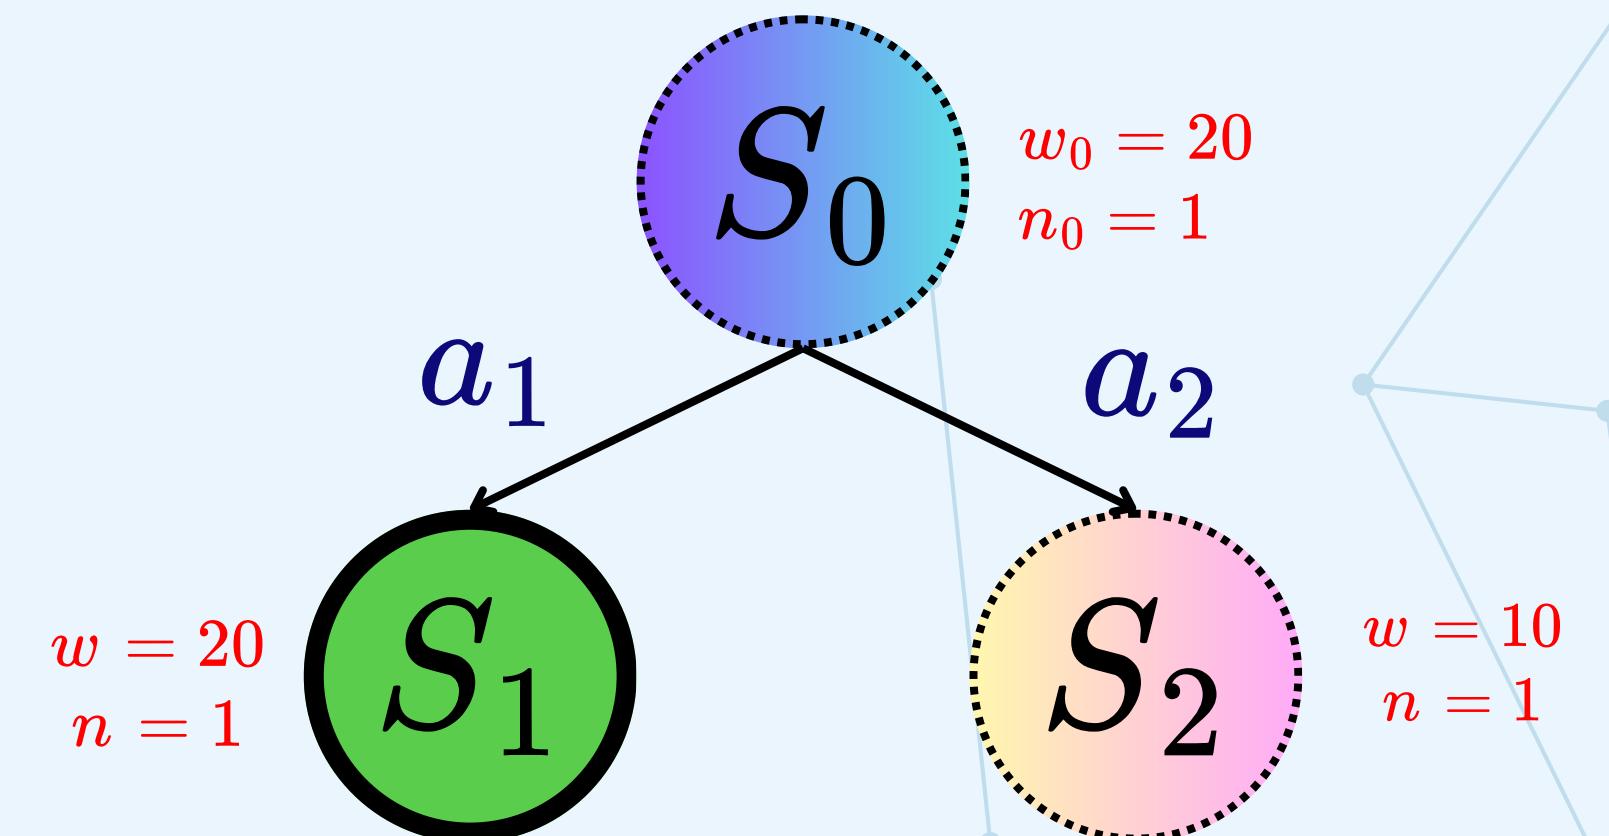
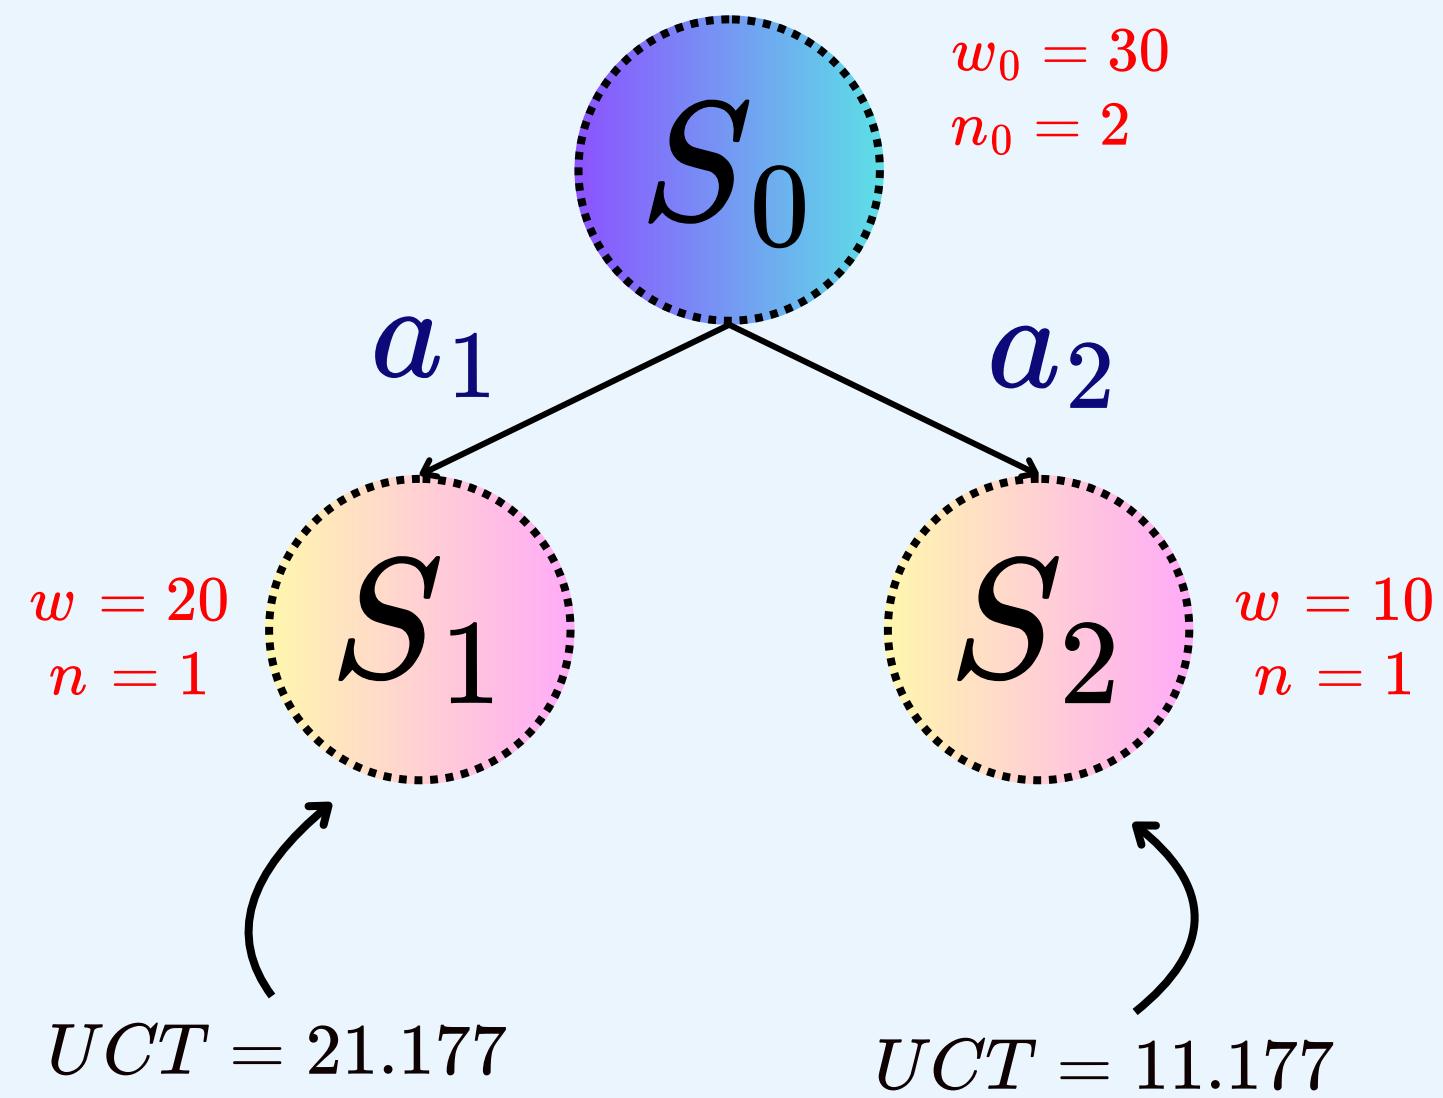


**Select  $S_2$  for second iteration**

## Second Iteration



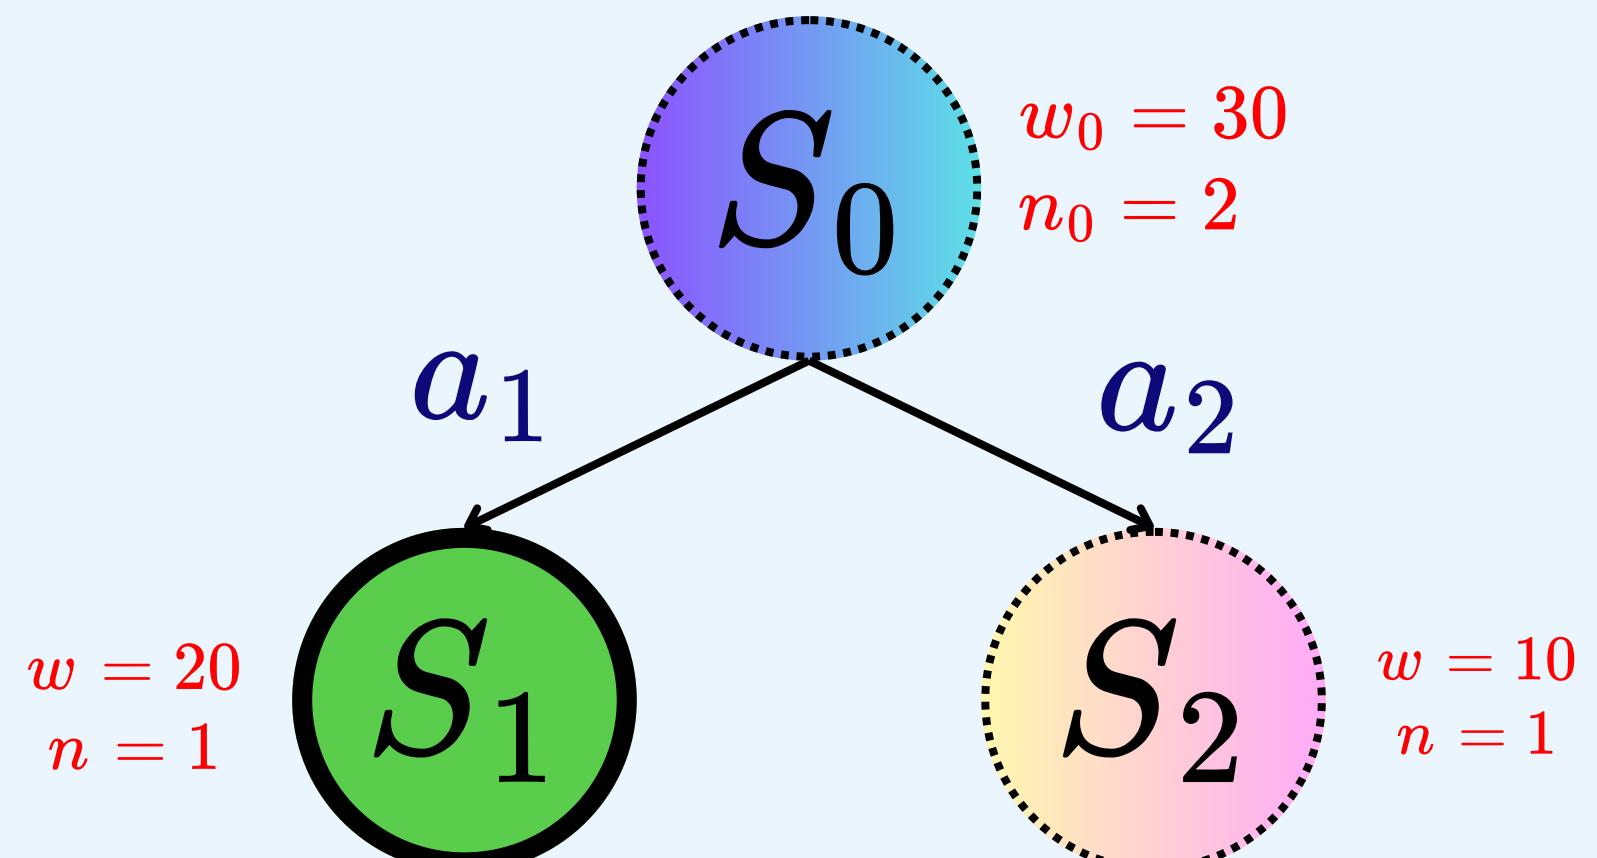
# Third Iteration



**Select  $S_1$  for third iteration**

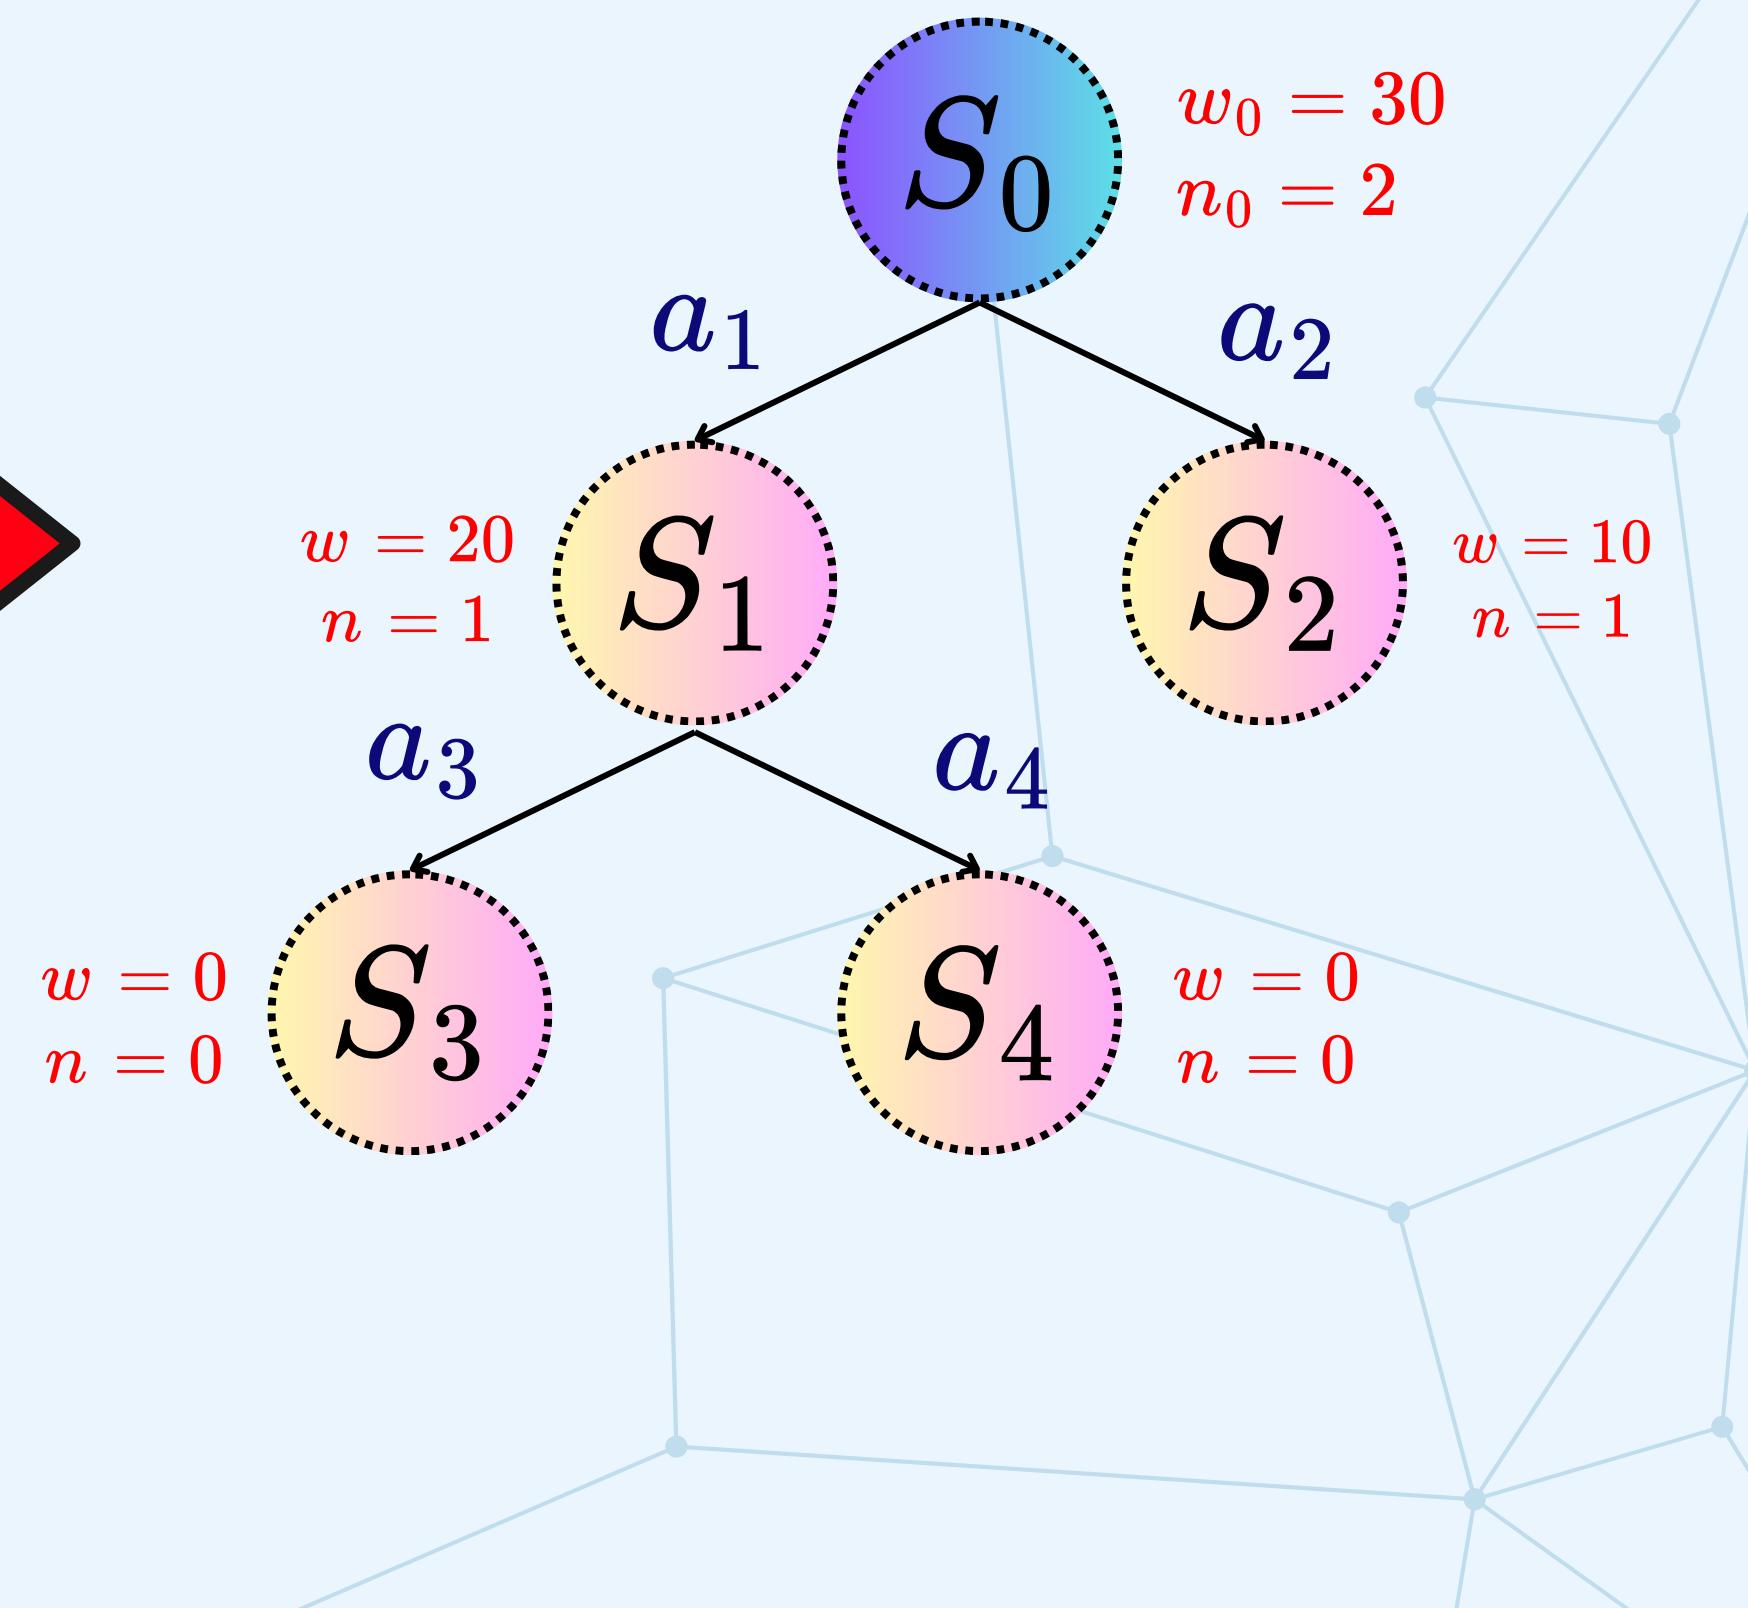
$$UCT = \frac{w_i}{n_i} + c \sqrt{\frac{\ln N_i}{n_i}}$$

# Third Iteration

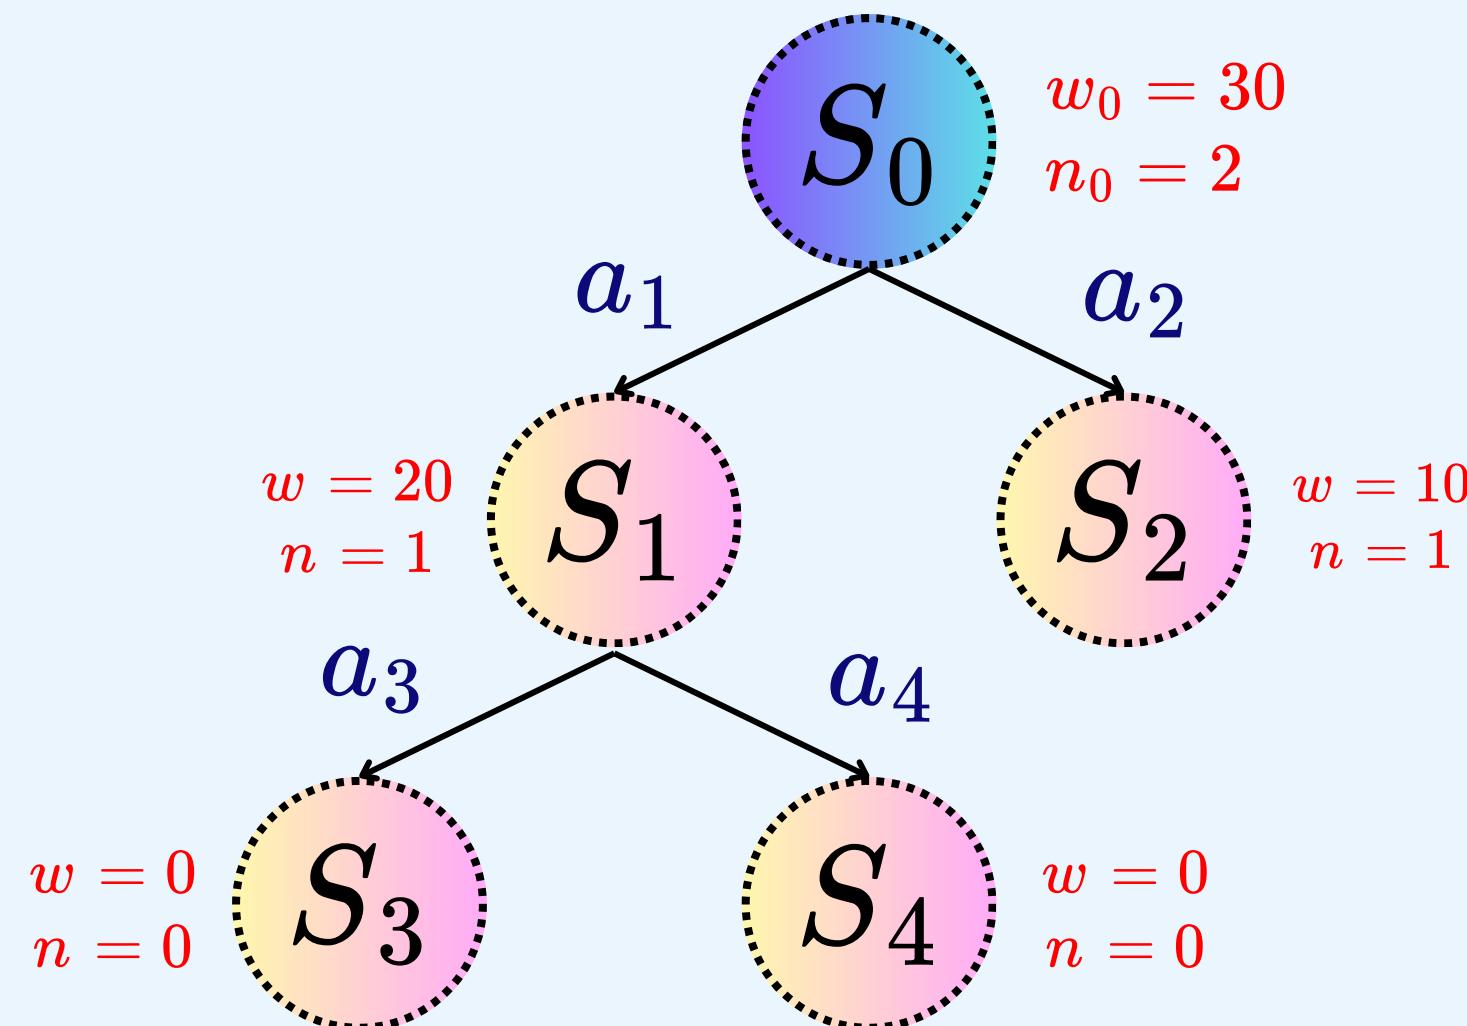


**EXPAND !**

Since  $S_1$  has been visited before

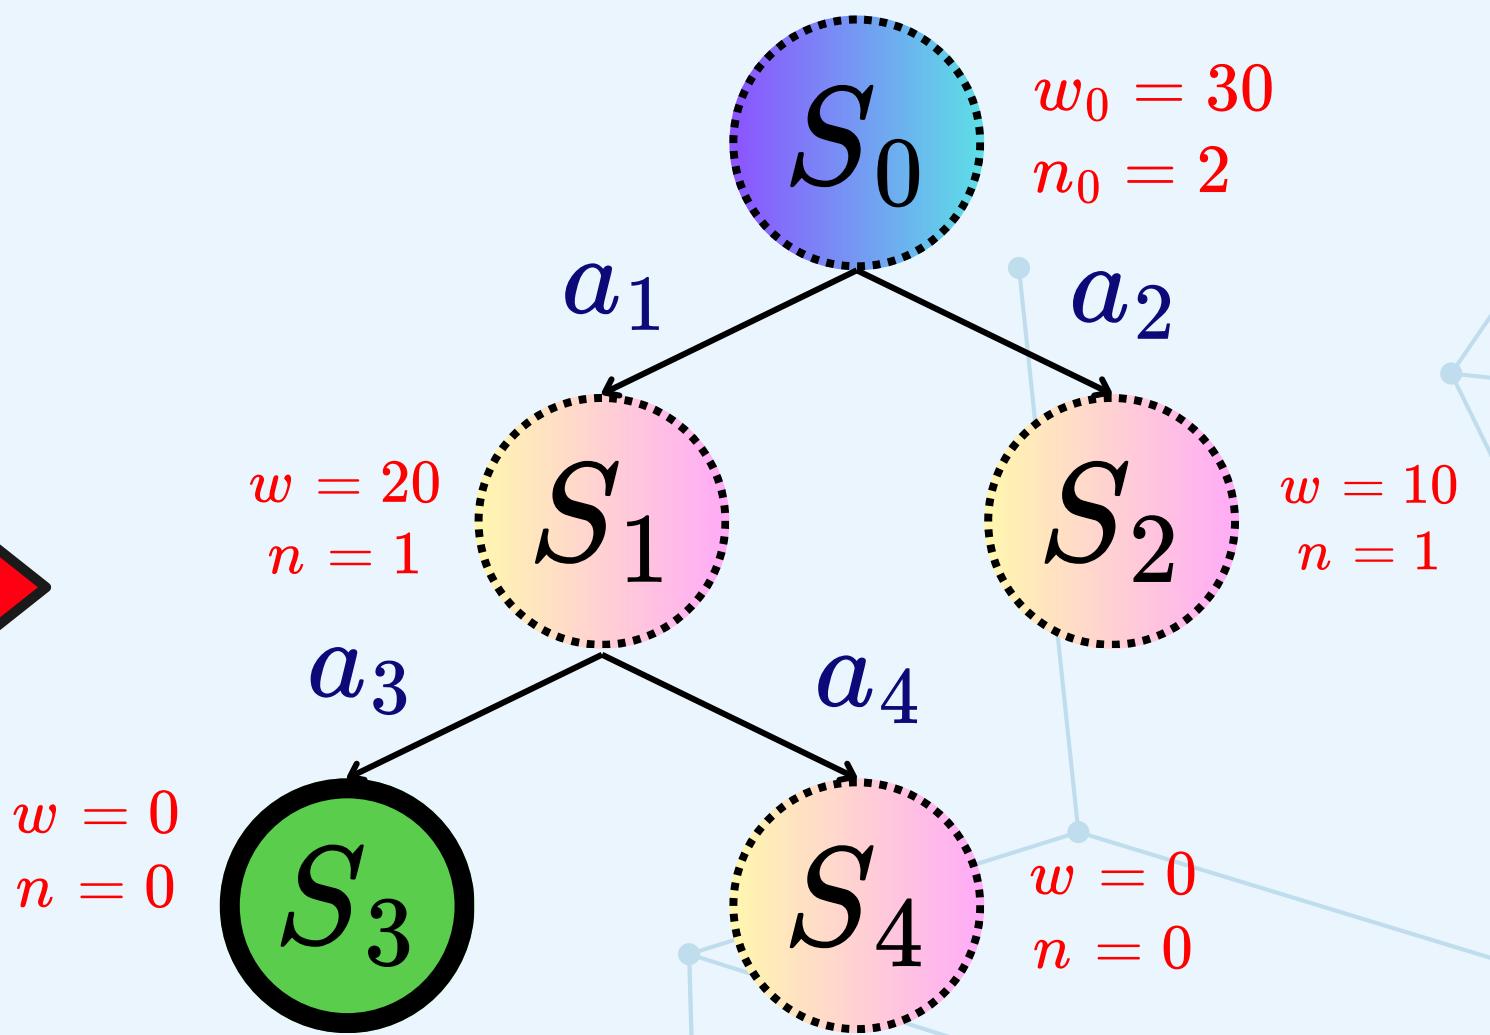


# Third Iteration



$UCT = \infty$

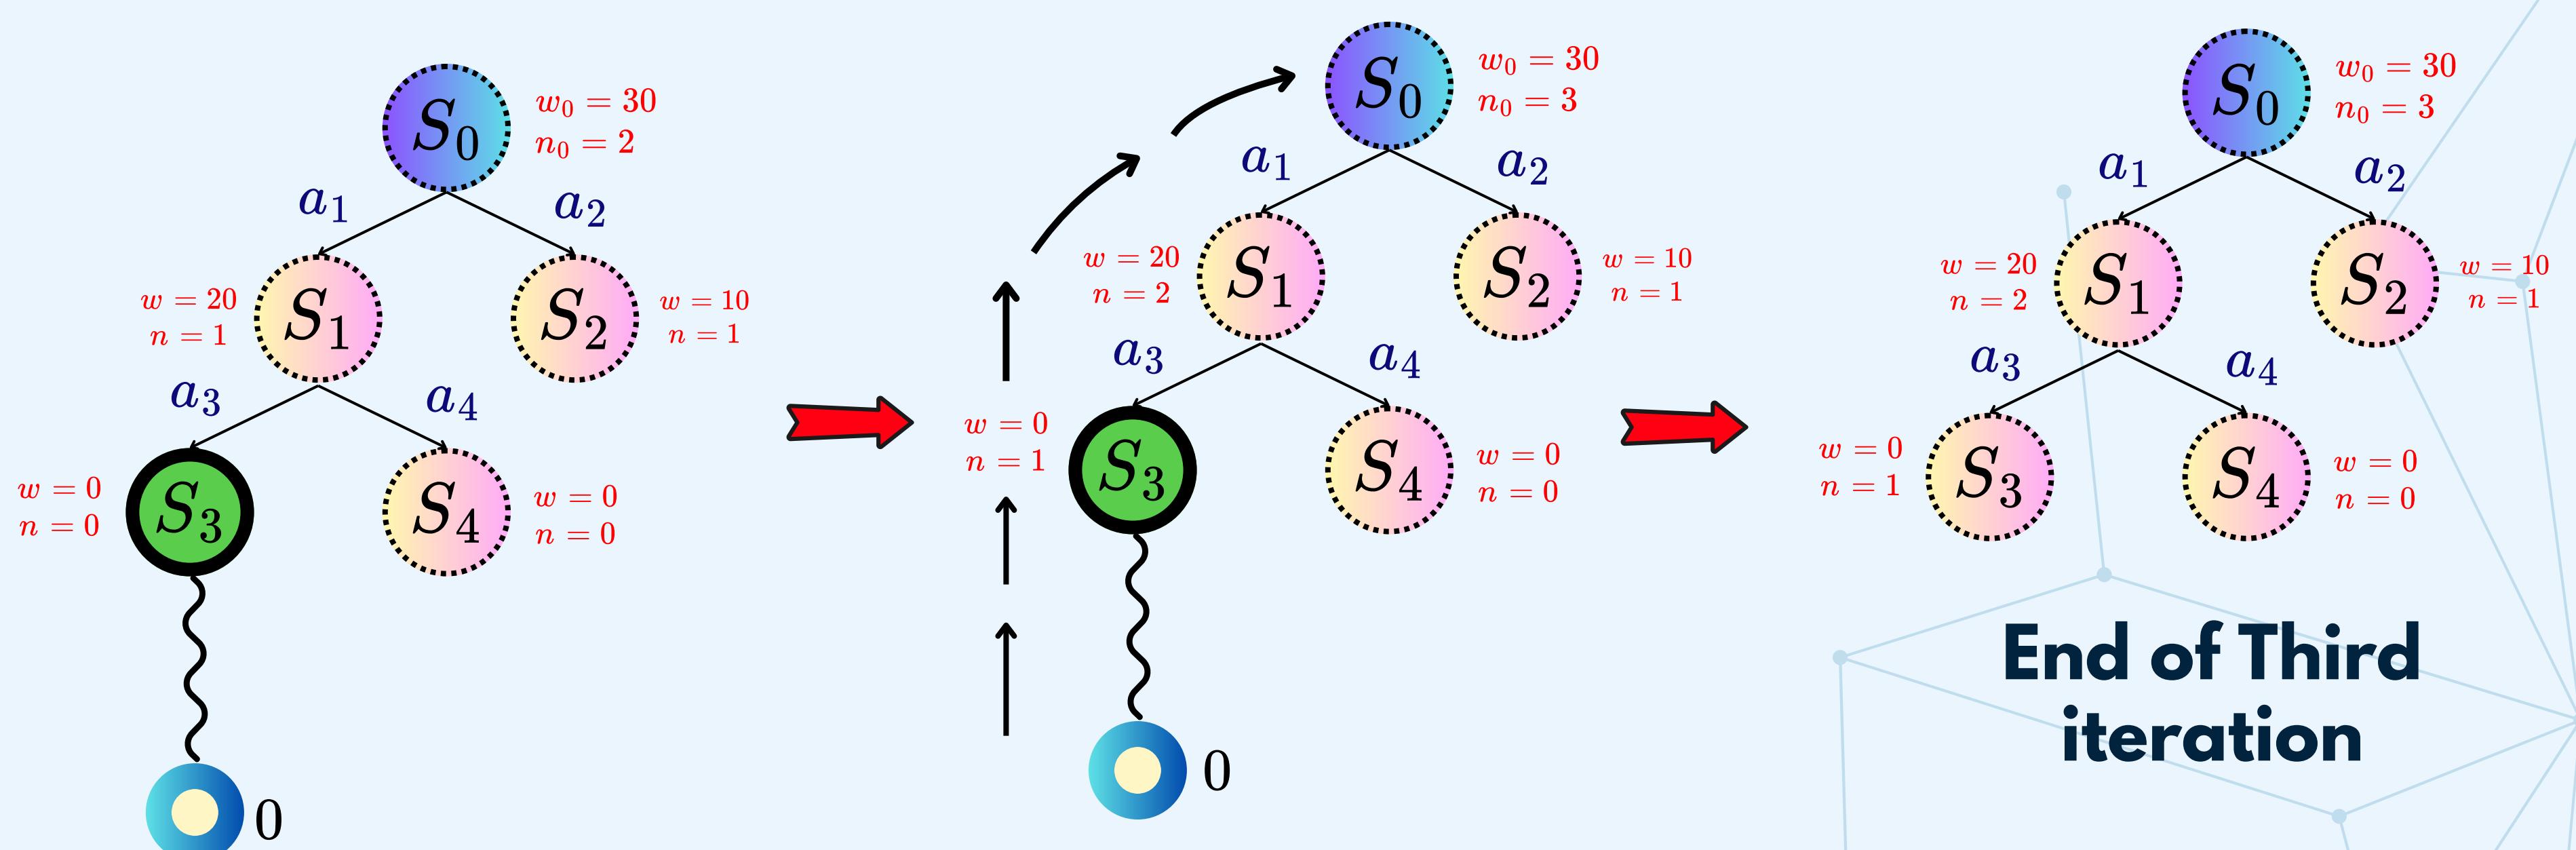
CHOOSE RANDOM



Since  $S_1$  has not been visited yet

ROLLOUT !

# Third Iteration



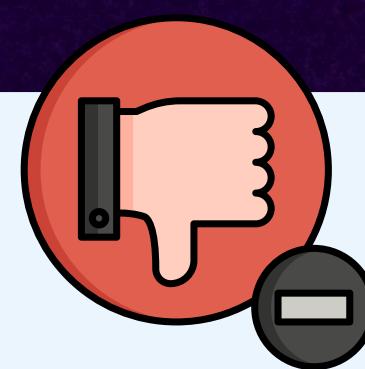
**End of Third  
iteration**

# Advantages and Disadvantages



## Advantages:

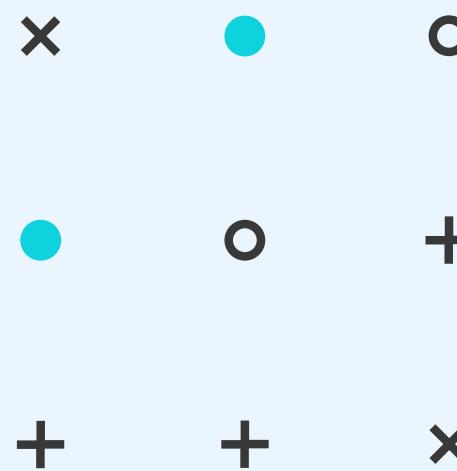
- Does not require a heuristic evaluation function
- Works well in large or complex state spaces
- Dynamically balances exploration and exploitation



## Disadvantages:

- Slow if you don't have enough time for many simulations
- Randomness makes it inconsistent unless repeated many times
- May miss deep tactical threats if simulation depth is shallow





# Common Q&A



## Q1: What's the difference between MCTS and Minimax?

- Minimax explores the full tree with evaluation at leaves
- MCTS explores promising paths using simulations, no evaluation function needed



## Q2: Can MCTS be used in non-game scenarios?

- MCTS also used in:
- Robotics
  - Network path optimization
  - Resource allocation



## Q3: How many iterations should we run?

As many as your time allows! The more simulations, the better the result.  
or when are sure about what is the best action



# THANK YOU!



VRQUEST

VR